

IMPLEMENTING OPERATIONS IN POWER-OF-2 CYCLOTOMIC RINGS

LATTICE MEETING

Martin R. Albrecht

2016-04-14

OUTLINE

GGH-like Multilinear Maps

Multiplication

Computing Algebraic Norms

Primality

Inverting in $\mathbb{Q}[X]/(X^n + 1)$

Small Remainders

Discrete Gaussians

Approximate Square Roots

OUTLINE

GGH-like Multilinear Maps

Multiplication

Computing Algebraic Norms

Primality

Inverting in $\mathbb{Q}[X]/(X^n + 1)$

Small Remainders

Discrete Gaussians

Approximate Square Roots

GGH-LIKE MULTILINEAR MAPS

- In 2013, Garg, Gentry and Halevi¹ proposed a construction, relying on ideal lattices, of a graded encoding scheme that approximates a cryptographic multilinear map.
- Shortly after, this construction was improved by Langlois, Stéhlé and Steinfeld².
- Implementing GGH-like schemes naively would not allow instantiating it for non-trivial parameter sizes.

¹Sanjam Garg, Craig Gentry, and Shai Halevi. [Candidate Multilinear Maps from Ideal Lattices](#). In: *EUROCRYPT 2013*. Ed. by Thomas Johansson and Phong Q. Nguyen. Vol. 7881. LNCS. Springer, Heidelberg, May 2013, pp. 1–17. DOI: 10.1007/978-3-642-38348-9_1.

²Adeline Langlois, Damien Stehlé, and Ron Steinfeld. [GGHlite: More Efficient Multilinear Maps from Ideal Lattices](#). In: *EUROCRYPT 2014*. Ed. by Phong Q. Nguyen and Elisabeth Oswald. Vol. 8441. LNCS. Springer, Heidelberg, May 2014, pp. 239–256. DOI: 10.1007/978-3-642-55220-5_14.

Martin R. Albrecht, Catalin Cocis, Fabien Laguillaumie, and Adeline Langlois. **Implementing Candidate Graded Encoding Schemes from Ideal Lattices**. In: *ASIACRYPT 2015, Part II*. ed. by Tetsu Iwata and Jung Hee Cheon. Vol. 9453. LNCS. Springer, Heidelberg, 2015, pp. 752–775. DOI: [10.1007/978-3-662-48800-3_31](https://doi.org/10.1007/978-3-662-48800-3_31)

WAIT, AREN'T THOSE ALL BROKEN?



<http://malb.io/are-graded-encoding-schemes-broken-yet.html>

Key Exchange

Yupu Hu and Huiwen Jia. [Cryptanalysis of GGH Map](#). accepted at EUROCRYPT 2016. 2015

- Polynomial-time attack using low-level encodings of zero.³

³Sage implementation: <https://martinralfbrecht.wordpress.com/2015/04/13/>

Attacks without Low-Level Encodings of Zero

Jung Hee Cheon, Jinhyuck Jeong, and Changmin Lee. *An Algorithm for NTRU Problems and Cryptanalysis of the GGH Multilinear Map without an encoding of zero*. In: *IACR Cryptology ePrint Archive* 2016 (2016). URL: <http://ia.cr/2016/139>

Martin Albrecht, Shi Bai, and Léo Ducas. *A subfield lattice attack on overstretched NTRU assumptions: Cryptanalysis of some FHE and Graded Encoding Schemes*. In: *IACR Cryptology ePrint Archive* 2016 (2016). URL: <http://ia.cr/2016/127>

- Polynomial-time attack for large levels of multilinearity κ without low-level encodings of zero.
- Subexponential attack for large levels of multilinearity κ without low-level encodings of zero without using the zero-testing parameter.

Indistinguishability Obfuscation

Eric Miles, Amit Sahai, and Mark Zhandry. *Annihilation Attacks for Multilinear Maps: Cryptanalysis of Indistinguishability Obfuscation over GGH13*. In: *IACR Cryptology ePrint Archive 2016* (2016). URL: <http://ia.cr/2016/147>

- Polynomial-time attack on several iO constructions⁴⁵⁶

⁴Sanjam Garg et al. *Candidate Indistinguishability Obfuscation and Functional Encryption for all Circuits*. In: *54th FOCS*. IEEE Computer Society Press, Oct. 2013, pp. 40–49.

⁵Boaz Barak et al. *Protecting Obfuscation against Algebraic Attacks*. In: *EUROCRYPT 2014*. Ed. by Phong Q. Nguyen and Elisabeth Oswald. Vol. 8441. LNCS. Springer, Heidelberg, May 2014, pp. 221–238. DOI: 10.1007/978-3-642-55220-5_13.

⁶Prabhanjan Vijendra Ananth, Divya Gupta, Yuval Ishai, and Amit Sahai. *Optimizing Obfuscation: Avoiding Barrington’s Theorem*. In: *ACM CCS 14*. Ed. by Gail-Joon Ahn, Moti Yung, and Ninghui Li. ACM Press, Nov. 2014, pp. 646–658.

WHAT GOOD IS AN IMPLEMENTATION OF GGH?

GGH-like graded encodings schemes might be broken, but designers of lattice-based schemes might still be tempted to write: “Sample $g \leftarrow D_{R,\sigma}$ until $\mathcal{I} = (g)$ is a prime ideal” or “Sample $f \leftarrow D_{(g)+c,\sigma}$.”

- We work in the m -th cyclotomic ring for m a power of two.
- It has degree $n = m/2$ and we consider the representation $R \simeq \mathbb{Z}[X]/(x^n + 1)$.
- We also consider $R_q \simeq \mathbb{Z}_q[X]/(x^n + 1)$ and $R_g \simeq \mathbb{Z}[X]/(x^n + 1, g)$.

GGHLITE: INSTANCE GENERATION

- **Instance generation.** Given security parameter λ and multilinearity parameter κ , determine scheme parameters $n, q, \sigma, \sigma', \ell_{g^{-1}}, \ell_b, \ell$ as in GGHLite⁷. Then proceed as follows:
 - Sample $g \leftarrow D_{R, \sigma}$ until $\|g^{-1}\| \leq \ell_{g^{-1}}$ and $\mathcal{I} = (g)$ is a **prime ideal**. Define encoding domain $R_g = R / (g)$.
 - Sample $z_i \leftarrow U(R_q)$ for all $0 < i \leq \kappa$.
 - Sample $h \leftarrow D_{R, \sqrt{q}}$ s.t. h and g are **co-prime** and define the zero-testing parameter $p_{zt} = \left[\frac{h}{g} \prod_{i=1}^{\kappa} z_i \right]_q$.
 - Return public parameters $\text{params} = (n, q, \ell)$ and p_{zt} .

⁷Adeline Langlois, Damien Stehlé, and Ron Steinfeld. **GGHLite: More Efficient Multilinear Maps from Ideal Lattices**. In: *EUROCRYPT 2014*. Ed. by Phong Q. Nguyen and Elisabeth Oswald. Vol. 8441. LNCS. Springer, Heidelberg, May 2014, pp. 239–256. DOI: 10.1007/978-3-642-55220-5_14.

- **Encode at level-0.** Compute a **small representative** $e' = [e]_g$ and **sample an element** $e'' \leftarrow D_{e'+\mathcal{I}, \sigma'}$. Output e'' .
- **Encode in group.** Given parameters params , z_i and a level-0 encoding $e \in R$, output $[e/z_i]_q$.

GGHLITE: ARITHMETIC & ZERO-TESTING

- **Adding encodings.** Given encodings $u_1 = [c_1 / (\prod_{i \in S} z_i)]_q$ and $u_2 = [c_2 / (\prod_{i \in S} z_i)]_q$ with $S \subseteq \{1, \dots, \kappa\}$:
 - Return $u = [u_1 + u_2]_q$, an encoding of $[c_1 + c_2]_q$ in the group S .
- **Multiplying encodings.** Let $S_1 \subset [\kappa]$, $S_2 \subset [\kappa]$ with $S_1 \cap S_2 = \emptyset$, given an encoding $u_1 = [c_1 / (\prod_{i \in S_1} z_i)]_q$ and an encoding $u_2 = [c_2 / (\prod_{i \in S_2} z_i)]_q$:
 - Return $u = [u_1 u_2]_q$, an encoding of $[c_1 c_2]_q$ in $S_1 \cup S_2$.
- **Zero testing.** Given parameters params , a zero-testing parameter p_{zt} , and an encoding $u = [c / (\prod_{i=0}^{\kappa-1} z_i)]_q$ in the group $[\kappa]$, return 1 if $\|[p_{zt} u]_q\|_\infty < q^{3/4}$ and 0 else.

```
gghlite-flint
|- applications    # 0.9k benchmarks, high-level applications, ...
|- dgs             # 1.2k discrete Gaussian sampling over the Integers
|- dgs1           # 0.7k discrete Gaussian sampling over lattices
|- flint           # 250k we rely on flint
|- gghlite         # 1.5k instance generation, zero testing ...
|- oz             # 2.2k operations in  $\mathbb{Z}[x]/(x^{n+1})$ 
|- tests           # 1.1k tests!
```

<https://bitbucket.org/malb/gghlite-flint>

<https://bitbucket.org/malb/dgs>

OUTLINE

GGH-like Multilinear Maps

Multiplication

Computing Algebraic Norms

Primality

Inverting in $\mathbb{Q}[X]/(X^n + 1)$

Small Remainders

Discrete Gaussians

Approximate Square Roots

OPTIONS

- Naive multiplication takes $\mathcal{O}(n^2)$.
- Asymptotically fast multiplication:
 - Reduce to multiplication in $\mathbb{Z}[X]$
 - Schönhage-Strassen algorithm for multiplying large integers in $\mathcal{O}(n \log n \log \log n)$.
 - This is the strategy implemented in FLINT.
 - FLINT has highly optimised implementation of the Schönhage-Strassen algorithm.
- We can also achieve $\mathcal{O}(n \log n)$ by the Number-Theoretic Transform.

NEGATIVE WRAPPED CONVOLUTION

Theorem (Negative Wrapped Convolution)

Let ω_n be an n th root of unity in \mathbb{Z}_q and $\varphi^2 = \omega_n$. Let

$$a = \sum_{i=0}^{n-1} a_i X^i \text{ and } b = \sum_{i=0}^{n-1} b_i X^i \in \mathbb{Z}_q[X]/(X^n + 1).$$

Let $c = a \cdot b \in \mathbb{Z}_q[X]/(X^n + 1)$ and let

$$\bar{a} = (a_0, \varphi a_1, \dots, \varphi^{n-1} a_{n-1})$$

and define \bar{b} and \bar{c} analogously. Then

$$\bar{c} = 1/n \cdot \text{NTT}_{\omega_n}^{-1}(\text{NTT}_{\omega_n}(\bar{a}) \odot \text{NTT}_{\omega_n}(\bar{b})).$$

NTT OVER MACHINE WORDS: BIT REVERSAL

```
void _nmod_vec_oz_ntt(mp_ptr rop, const mp_ptr op, const mp_ptr w,
                     const size_t n, const nmod_t q) {
    const size_t k = n_flog(n,2);

    mp_ptr a = _nmod_vec_init(n);

    for (unsigned int i = 0; i < n; i++) {
        unsigned int r;
        r = (bit_reverse_table_256[i & 0xff] << 16) | \
            (bit_reverse_table_256[(i >> 8) & 0xff] << 8) | \
            (bit_reverse_table_256[(i >> 16) & 0xff]);
        r >>= (24 - k);
        a[r] = op[i];
    }
}
```

NTT OVER MACHINE WORDS: MAIN LOOP

```
mp_ptr b = _nmod_vec_init(n);

const double ninv = n_precompute_inverse(q.n);
for(size_t i=0; i<k; i++) {
    const mp_limb_t tkm = ~(((1UL)<<(k-1-i)) - 1);
    for(size_t j=0; j<n/2; j++) {
        const size_t pij = j & tkm;
        mp_limb_t tmp = n_mulmod_precomp(a[2*j+1], w[pij], q.n, ninv);
        b[j]          = n_addmod(a[2*j], tmp, q.n);
        b[j+n/2]      = n_submod(a[2*j], tmp, q.n);
    }
    if(i!=k-1)
        _nmod_vec_set(a, b, n);
}
_nmod_vec_set(rop, b, n);
_nmod_vec_clear(b);
_nmod_vec_clear(a);
}
```

AVOIDING CONVERSION

- If we do many operations in $\mathbb{Z}_q[X]/(X^n + 1)$ we can avoid repeated conversions between coefficient and “evaluation” representation $(f(1), f(\omega_n), \dots, f(\omega_n^{n-1}))$
- We convert encodings to their evaluation representation once on creation
- We convert back only when running extraction.
- This reduces the amortised cost from $\mathcal{O}(n \log n)$ to $\mathcal{O}(n)$.

OUTLINE

GGH-like Multilinear Maps

Multiplication

Computing Algebraic Norms

Primality

Inverting in $\mathbb{Q}[X]/(X^n + 1)$

Small Remainders

Discrete Gaussians

Approximate Square Roots

COMPUTING ALGEBRAIC NORMS: RESULTANTS

- During instance generation we have to compute the norm of g .
- We can compute norms in $\mathbb{Z}[X]/(X^n + 1)$ by observing that

$$\mathcal{N}(f) = \text{res}(f, X^n + 1).$$

MULTI-MODULAR RESULTANTS

- The usual strategy for computing resultants over the integers is to use a multi-modular approach.
- We compute resultants modulo many small primes q_i and then combine the results using the Chinese Remainder Theorem.
- Resultants modulo a prime q_i can be computed in $\mathcal{O}(M(n) \log n)$ operations where $M(n)$ is the cost of one multiplication in $\mathbb{Z}_{q_i}[X]/(X^n + 1)$.
- Overall cost $\mathcal{O}(n \log^2 n)$ without specialisation.

- $\text{res}(f, X^n + 1) \bmod q_i$ can be rewritten as

$$\prod_{(X^n+1)(x)=0} f(x) \bmod q_i,$$

i.e. as evaluating f on all roots of $X^n + 1$.

- Picking q_i such that $q_i \equiv 1 \bmod 2n$ this can be accomplished using the NTT reducing the cost mod q_i to $\mathcal{O}(M(n))$ saving a factor of $\log n$.

SOURCE CODE: MAIN LOOP

```
void _fmpz_poly_oz_ideal_norm(fmpz_t norm, const fmpz_poly_t f,
                             const long n) {
    ...
    #pragma omp parallel for
    for (i = 0; i < num_primes; i++) {
        nmod_t mod;
        nmod_init(&mod, parr[i]);

        const int id = omp_get_thread_num();
        /* reduce polynomials modulo p */
        _fmpz_vec_get_nmod_vec(a[id], F, n, mod);
        /* compute resultant over Z/pZ */
        rarr[i] = _nmod_vec_oz_resultant(a[id], n, mod);
        flint_cleanup();
    }
    ...
}
```

SOURCE CODE: RESULTANTS MOD $q \equiv 1 \pmod{2n}$

```
mp_limb_t _nmod_vec_oz_resultant(const mp_ptr a, long n, nmod_t q) {
    const mp_limb_t w_ = _nmod_nth_root(2*n, q.n);
    mp_ptr w = _nmod_vec_init(2*n);
    mp_ptr t = _nmod_vec_init(2*n);

    _nmod_vec_oz_set_powers(w, 2*n, w_, q);
    _nmod_vec_oz_ntt(t, a, w, 2*n, q);

    mp_limb_t acc = 1;
    for(int i=1; i<2*n; i+=2)
        acc = n_mulmod2_preinv(acc, t[i], q.n, q.ninv);

    _nmod_vec_clear(w);
    _nmod_vec_clear(t);
    return acc;
}
```

OUTLINE

GGH-like Multilinear Maps

Multiplication

Computing Algebraic Norms

Primality

Inverting in $\mathbb{Q}[X]/(X^n + 1)$

Small Remainders

Discrete Gaussians

Approximate Square Roots

CHECKING PRIMALITY

To check if (g) is prime, compute the norm and check if prime. This is a sufficient but not necessary condition.

```
int fmpz_poly_ideal_is_probaprime(const fmpz_poly_t f,
                                  const long n,
                                  const mp_limb_t *primes) {
    ...
    fmpz_t norm;
    fmpz_init(norm);
    fmpz_poly_ideal_norm(norm, f, n, 0);
    r = fmpz_is_probabprime(norm);
    fmpz_clear(norm);
    ...
    return r;
}
```

RULING OUT COMMON FACTORS QUICKLY

Before computing resultants, check if $\text{res}(g, X^n + 1) \equiv 0 \pmod{q_i}$ for several "interesting" primes q_i .

```
int fmpz_poly_oz_ideal_is_probaprime(const fmpz_poly_t f,
                                     const long n,
                                     const mp_limb_t *primes) {
    int r = fmpz_poly_oz_ideal_not_prime_factors(f, n, primes);
    if (r) {
        fmpz_t norm;
        fmpz_init(norm);
        fmpz_poly_oz_ideal_norm(norm, f, n, 0);
        r = fmpz_is_probabprime(norm);
        fmpz_clear(norm);
    }
    return r;
}
```

RULING OUT COMMON FACTORS QUICKLY

```
int fmpz_poly_oz_ideal_not_prime_factors(const fmpz_poly_t f, long n,
                                         const mp_limb_t *primes) {
    nmod_poly_t a[num_threads], b[num_threads];
    int r[num_threads];

    for(size_t i=0; i<k; i+=num_threads) {
        if (k-i < (unsigned long)num_threads)
            num_threads = k-i;
    #pragma omp parallel for
        for (int j=0; j<num_threads; j++) {
            mp_limb_t p = primes[1+i+j];
            fmpz_poly_get_nmod_poly(a[j], f);
            r[j] = nmod_poly_oz_resultant(a[j], n);
        }
        for(int j=0; j<num_threads; j++)
            if (r[j] == 0)
                return r[0];
    }
    return r[0];
}
```

COMMON FACTORS

These primes are 2 and then all primes up to some bound with $q_i \equiv 1 \pmod n$ because these occur with good probability as factors.

```
int _gghlite_nsmall_primes(const gghlite_params_t self) {  
    /* we try about 1% small primes first, where 1% relates to the total  
       number of primes needed for multi-modular result */  
    const long n = self->n;  
    int nsp = ceil((log2(_gghlite_sigma(n)) + log2(n)/2.0)  
                  * n/100.0/(FLINT_BITS -1));  
    if (nsp < 20)  
        nsp = 20;  
    return nsp;  
}
```


TIMINGS

n	$\log \sigma$	wall time
1024	15.1	0.54s
2048	16.2	3.03s
4096	17.3	20.99s
32768	20.4	1834.99s

Average time of checking primality of a single (g) on Intel Xeon CPU E5-2667 v2 3.30GHz with 256GB of RAM using 16 cores.

VERIFYING CO-PRIMALITY

- When re-randomisation elements are required, then it is necessary that they generate all of (g) , i.e.

$$(b_1^{(1)}, b_2^{(1)}) = (g).$$

- When $b_i^{(1)} = \tilde{b}_i^{(1)}g$ for $0 < i \leq 2$ then this is equivalent to

$$(\tilde{b}_1^{(1)}) + (\tilde{b}_2^{(1)}) = \mathbb{Z}[X]/(X^n + 1).$$

- We check the sufficient but not necessary condition

$$\gcd(\text{res}(\tilde{b}_1^{(1)}, X^n + 1), \text{res}(\tilde{b}_2^{(1)}, X^n + 1)) = 1,$$

i.e. if the respective ideal norms are co-prime.

AVOIDING RESULTANTS

- Perform this check for every candidate pair $(\tilde{b}_1^{(1)}, \tilde{b}_2^{(1)})$.
- Compute two resultants and their gcd: **expensive**.
- But

$$\gcd(\text{res}(\tilde{b}_1^{(1)}, X^n + 1), \text{res}(\tilde{b}_2^{(1)}, X^n + 1)) \neq 1$$

when

$$\text{res}(\tilde{b}_1^{(1)}, X^n + 1) = 0 = \text{res}(\tilde{b}_2^{(1)}, X^n + 1) \bmod q_i$$

for any modulus q_i .

➔ Check this condition for several “interesting” primes and resample if this condition holds.

AVOIDING RESULTANTS

- After having ruled out small common prime factors it is quite unlikely that the gcd of the norms is not equal to one.
- With good probability we will perform this expensive step only once as a final verification.

Improvement

A possible strategy is to sample $m > 2$ re-randomisers $b_i^{(1)}$ and to apply bounds on the probability of m elements $\tilde{b}_i^{(1)}$ sharing a prime factor after excluding small prime factors.

OUTLINE

GGH-like Multilinear Maps

Multiplication

Computing Algebraic Norms

Primality

Inverting in $\mathbb{Q}[X]/(X^n + 1)$

Small Remainders

Discrete Gaussians

Approximate Square Roots

Instance generation relies on inversion in $\mathbb{Q}[X]/(X^n + 1)$.

1. when sampling g we have to check that the norm of its inverse is bounded by ℓ_g .
2. To set up our discrete Gaussian samplers we need to run many inversions in an iterative process.

INVERTING IN $\mathbb{Q}[X]/(X^n + 1)$

- The core idea⁸ is similar to the FFT, i.e. to reduce the inversion of f to the inversion of an element of degree $n/2$.
- Since n is even, $f(X)$ is invertible modulo $X^n + 1$ if and only if $f(-X)$ is also invertible.
- By setting

$$F(X^2) = f(X)f(-X) \bmod X^n + 1,$$

the inverse $f^{-1}(X)$ of $f(X)$ satisfies

$$F(X^2)f^{-1}(X) = f(-X) \bmod X^n + 1.$$

⁸Dario Bini, Gianna M. Del Corso, Giovanni Manzini, and Luciano Margara. *Inversion of Circulant Matrices over \mathbb{Z}_m* . In: *International Colloquium on Automata, Languages and Programming*. Vol. 1443. LNCS. Springer, 1998, pp. 719–730.

INVERTING IN $\mathbb{Q}[X]/(X^n + 1)$

- Let

$$f^{-1}(X) = g(X) = G_e(X^2) + XG_o(X^2)$$

and

$$f(-X) = F_e(X^2) + XF_o(X^2).$$

- We obtain

$$F(X^2)(G_e(X^2) + XG_o(X^2)) = F_e(X^2) + XF_o(X^2) \bmod X^n + 1$$

or equivalently

$$F(X^2)G_e(X^2) = F_e(X^2) \pmod{X^n + 1},$$

$$F(X^2)G_o(X^2) = F_o(X^2) \pmod{X^n + 1}$$

- Invert $f(X)$ by inverting $F(X^2)$ and multiplying at degree $n/2$.
- Recursively call the inversion of $F(Y)$ modulo $(X^{n/2} + 1)$ by setting $Y = X^2$.

Instance generation relies on inversion in $\mathbb{Q}[X]/(X^n + 1)$.

1. when sampling g we have to check that the norm of its inverse is bounded by ℓ_g .
2. To set up our discrete Gaussian samplers we need to run many inversions in an iterative process.

GGH MOTIVATION REVISITED

Instance generation relies on inversion in $\mathbb{Q}[X]/(X^n + 1)$.

1. when sampling g we have to check that the norm of its inverse is bounded by ℓ_g .
2. To set up our discrete Gaussian samplers we need to run many inversions in an iterative process.

Approximates Suffice

In the first case we only need to estimate the size of g^{-1} and in the second case inversion is a subroutine of an approximation algorithm.

TRUNCATION

```
void fmpq_poly_truncate_prec(fmpq_poly_t op, const mp_bitcnt_t prec) {
    mpq_t *tmp_q = (mpq_t*)calloc(fmpq_poly_length(op), sizeof(mpq_t));
    mpf_t tmp_f; mpf_init2(tmp_f, prec);

    for (int i=0; i<fmpq_poly_length(op); i++) {
        mpq_init(tmp_q[i]);
        fmpq_poly_get_coeff_mpq(tmp_q[i], op, i);
        mpf_set_q(tmp_f, tmp_q[i]);
        mpq_set_f(tmp_q[i], tmp_f);
    }
    fmpq_poly_set_array_mpq(op, (const mpq_t*)tmp_q, fmpq_poly_length(op));
    ...
}
```

Calling `fmpq_poly_set_array_mpq` instead of setting each coefficient one-by-one avoids repeated GCD computations.

ALGORITHM

if $n = 1$ then

$$g_0 \leftarrow f_0^{-1}$$

else

$$F(X^2) \leftarrow f(X)f(-X) \bmod X^n + 1$$

$$\tilde{F}(Y) = F(Y) \text{ truncated to } \mathbf{prec} \text{ bits of precision}$$

$$G(Y) \leftarrow \text{InverseMod}(\tilde{F}(Y), q, n/2)$$

$$\text{Set } F_e(X^2), F_o(X^2) \text{ such that } f(-X) = F_e(X^2) + XF_o(X^2)$$

$$T_e(Y), T_o(Y) \leftarrow G(Y)F_e(Y), G(Y)F_o(Y)$$

$$f^{-1}(X) \leftarrow T_e(X^2) + XT_o(X^2)$$

$$\tilde{f}^{-1}(X) = f^{-1}(X) \text{ truncated to } \mathbf{prec} \text{ bits of precision}$$

$$\text{return } \tilde{f}^{-1}(X)$$

end if

Approximate inverse of $f(X) \bmod X^n + 1$ using **prec** bits of precision

TIMINGS

n	$\log \sigma$	xgcd	160	160iter	∞
4096	17.2	234.1s	0.067s	0.073s	121.8s
8192	18.3	1476.8s	0.195s	0.200s	755.8s

Inverting $g \leftarrow_{\$} D_{\mathbb{Z}^n, \sigma}$ with FLINT's extended Euclidean algorithm ("xgcd"), our implementation with precision 160 ("160"), iterating our implementation until $\|\tilde{f}^{-1}(X)f(X) - 1\| < 2^{-160}$ ("160iter") and our implementation without truncation (" ∞ ") on Intel Core i7-4850HQ CPU at 2.30GHz, single core.

OUTLINE

GGH-like Multilinear Maps

Multiplication

Computing Algebraic Norms

Primality

Inverting in $\mathbb{Q}[X]/(X^n + 1)$

Small Remainders

Discrete Gaussians

Approximate Square Roots

SMALL REMAINDERS: MOTIVATION

- The Jigsaw Generator⁹ takes as input elements a_i in \mathbb{Z}_p where $p = \mathcal{N}(\mathcal{I})$ and produces encodings with respect to some S_i .
- This algorithm produces some small representative of the coset a_i modulo (g) from large integers of size $\approx (\sigma\sqrt{n})^n$.

⁹Sanjam Garg et al. [Candidate Indistinguishability Obfuscation and Functional Encryption for all Circuits](#). In: *54th FOCS*. IEEE Computer Society Press, Oct. 2013, pp. 40–49.

- We can use Babai's trick and that g is small, i.e. compute

$$a_i - g \cdot \lfloor g^{-1} \cdot a_i \rfloor \text{ in } \mathbb{Q}[X]/(X^n + 1)$$

- To produce sufficiently small elements, we need g^{-1} either exactly or with high precision.
- Computing such a high quality approximation of g^{-1} is prohibitively expensive.

SMALL REMAINDERS: SD

1. Rewrite a_i as

$$a_i = \sum_{j=0}^{\lceil \log_2(a_i)/B \rceil} 2^{B \cdot j} \cdot a_{ij}$$

where $a_{ij} < 2^B$ for some B .

2. Compute small representatives for all $2^{B \cdot j}$ and a_{ij} using an approximation of g^{-1} with precision B .
3. Multiply small representatives for $2^{B \cdot j}$ and a_{ij} and add up their products.

This produces a somewhat short element which we then reduce using approximation of g^{-1} with precision B until its size does not decrease any more.

OUTLINE

GGH-like Multilinear Maps

Multiplication

Computing Algebraic Norms

Primality

Inverting in $\mathbb{Q}[X]/(X^n + 1)$

Small Remainders

Discrete Gaussians

Approximate Square Roots

DISCRETE GAUSSIAN SAMPLING

- We need to sample from the discrete Gaussian $D_{(g),\sigma',c}$ where c is a small representative of a coset of (g) .
- Fundamental building block is sampler over the Integers.

- Discrete Gaussian sampler over the integers for arbitrary precision using **MPFR** and **double** precision.
- Implements rejection sampling from a uniform distribution with and without table (“online”) lookups¹⁰ and Ducas et al’s sampler which samples from $D_{\mathbb{Z}, k\sigma_2}$ where σ_2 is a constant¹¹.
- Implementation automatically chooses the best algorithm based on σ , c and τ (tail cut).

¹⁰Craig Gentry, Chris Peikert, and Vinod Vaikuntanathan. [Trapdoors for hard lattices and new cryptographic constructions](#). In: *40th ACM STOC*. ed. by Richard E. Ladner and Cynthia Dwork. ACM Press, May 2008, pp. 197–206.

¹¹Léo Ducas, Alain Durmus, Tancrede Lepoint, and Vadim Lyubashevsky. [Lattice Signatures and Bimodal Gaussians](#). In: *CRYPTO 2013, Part I*. ed. by Ran Canetti and Juan A. Garay. Vol. 8042. LNCS. Springer, Heidelberg, Aug. 2013, pp. 40–56. DOI: [10.1007/978-3-642-40041-4_3](https://doi.org/10.1007/978-3-642-40041-4_3).

TIMINGS

algorithm	σ	c	prec	samp./s	prec	samp./s
tabulated	10000	1.0	53	660.000	160	310.000
tabulated	10000	0.5	53	650.000	160	260.000
online	10000	1.0	53	414.000	160	9.000
online	10000	0.5	53	414.000	160	9.000
Alg 12 [DDLL13]	10000	1.0	53	350.000	160	123.000

Example timings for discrete Gaussian sampling over \mathbb{Z} on Intel Core i7-4850HQ CPU at 2.30GHz, single core.

- Implemented naively this takes $\mathcal{O}(n^3 \log n)$ operations even if we ignore issues of precision.
- Following Léo's thesis¹², we implemented a variant of Peikert's sampler¹³.

¹²Léo Ducas. *Signatures Fondées sur les Réseaux Euclidiens: Attaques, Analyse et Optimisations*. PhD thesis. Université Paris Diderot, 2013.

¹³Chris Peikert. *An Efficient and Parallel Gaussian Sampler for Lattices*. In: *CRYPTO 2010*. Ed. by Tal Rabin. Vol. 6223. LNCS. Springer, Heidelberg, Aug. 2010, pp. 80–97.

SAMPLING FROM $D_{(g),\sigma',0}$

1. Observe that

$$D_{(g),\sigma',0} = g \cdot D_{R,\sigma'g^{-T}}$$

2. Compute approximate square-root $\sqrt[\text{appr}]{\Sigma_2}$ of

$$\Sigma_2 = \sigma'^2 \cdot g^{-T} \cdot g^{-1} - r^2 \text{ with } r = 2 \cdot \lceil \sqrt{\log n} \rceil$$

3. Sample a vector $x \leftarrow_{\$} \mathbb{R}^n$ from a standard normal distribution and interpret it as a polynomial in $\mathbb{Q}[X]/(X^n + 1)$.
4. Compute $y = \sqrt[\text{appr}]{\Sigma_2} \cdot x$ in $\mathbb{Q}[X]/(X^n + 1)$ and return $g \cdot (\lfloor y \rfloor_r)$, where $\lfloor y \rfloor_r$ denotes sampling a vector in \mathbb{Z}^n where the i -th component follows $D_{\mathbb{Z},r,y_i}$.

SAMPLING FROM $D_{(g),\sigma',0}$: SQRT

1. Compute an approximate square root of

$$\Sigma'_2 = g^{-T} \cdot g^{-1}$$

up to λ bits of precision.

- Precision: $\log(n) + 4(\log(\sqrt{n}\sigma))$ bits.
 - If square root does not converge, double precision and start over.
2. Use this approximate square-root, scaled appropriately, as the initial value from which to start computing a square-root of

$$\Sigma_2 = \sigma'^2 \cdot g^{-T} \cdot g^{-1} - r^2 \text{ with } r = 2 \cdot \lceil \sqrt{\log n} \rceil$$

3. Terminate when the square is within distance $2^{-2\lambda}$ to Σ_2 .
4. Converges quickly because initial candidate close to target.

OUTLINE

GGH-like Multilinear Maps

Multiplication

Computing Algebraic Norms

Primality

Inverting in $\mathbb{Q}[X]/(X^n + 1)$

Small Remainders

Discrete Gaussians

Approximate Square Roots

- For some input element Σ we want to compute some element $\sqrt[n]{\Sigma} \in \mathbb{Q}[X]/(X^n + 1)$ such that $\| \sqrt[n]{\Sigma} \cdot \sqrt[n]{\Sigma} - \Sigma \| < 2^{-2\lambda}$.
- We use iterative methods which iteratively refine the approximation of the square root similar to Newton's method.¹⁴
- Computing approximate square roots of matrices is a well studied research area with many algorithms known in the literature.¹⁵
- All algorithms with global convergence invoke approximate inversions in $\mathbb{Q}[X]/(X^n + 1)$ for which we call our inversion algorithm.

¹⁴Léo Ducas. *Signatures Fondées sur les Réseaux Euclidiens: Attaques, Analyse et Optimisations*. PhD thesis. Université Paris Diderot, 2013.

¹⁵Nicholas J. Higham. *Stable iterations for the matrix square root*. In: *Numerical Algorithms* 15.2 (1997), pp. 227–242. ISSN: 1017-1398. DOI: [10.1023/A:1019150005407](https://doi.org/10.1023/A:1019150005407). URL: <http://dx.doi.org/10.1023/A%3A1019150005407>.

Babylonian only one inversion, which allows lower precision.

Denman-Beavers converges faster in practice and can be parallelised on two cores.¹⁶

Padé iteration arbitrarily many cores, but workload on each core is greater than Denman-Beavers.¹⁷ Only better for us when more than 8 cores were used.

¹⁶Eugene D. Denman and Alex N. Beavers Jr. *The matrix sign function and computations in systems*. In: *Applied Mathematics and Computation* 2.1 (1976), pp. 63–94.

¹⁷Nicholas J. Higham. *Stable iterations for the matrix square root*. In: *Numerical Algorithms* 15.2 (1997), pp. 227–242. ISSN: 1017-1398. DOI: [10.1023/A:1019150005407](https://doi.org/10.1023/A:1019150005407). URL: <http://dx.doi.org/10.1023/A%3A1019150005407>.

RAPID CONVERGENCE

- Quadratic convergence does not assure rapid convergence in practice because error can take many iterations to become small enough.
- Speed-up convergence by scaling the operands appropriately in each loop.¹⁸
- Common scaling scheme: scale by the determinant, i.e. $\text{res}(f, X^n + 1)$ for some $f \in \mathbb{Q}[X] / (X^n + 1)$.
- Computing resultants in $\mathbb{Q}[X] / (X^n + 1)$ reduces to computing resultants in $\mathbb{Z}[X] / (X^n + 1)$.
- Computing resultants in $\mathbb{Z}[X] / (X^n + 1)$ can be expensive.

¹⁸Nicholas J. Higham. *Stable iterations for the matrix square root*. In: *Numerical Algorithms* 15.2 (1997), pp. 227–242. ISSN: 1017-1398. DOI: [10.1023/A:1019150005407](https://doi.org/10.1023/A:1019150005407). URL: <http://dx.doi.org/10.1023/A%3A1019150005407>.

APPROXIMATE RESULTANTS

- We are only interested in approximate determinant for scaling
→ compute with reduced precision.
- Clear all but the most significant bit for each coefficient's numerator and denominator of f to produce f' and compute $\text{res}(f', X^n + 1)$.
- Reduces the size of the integer representation to speed up the resultant computation.
- With this optimisation scaling by an approximation of the determinant is both fast and precise enough to produce fast convergence.

SQRT TIMING

prec	n	$\log \sigma'$	it.	wall time	$\log \left(\left(\sqrt[n]{\Sigma_2} \right)^2 - \Sigma_2 \right)$
160	1024	45.8	9	0.4s	-200
160	2048	49.6	9	0.9s	-221
160	4096	53.3	10	2.5s	-239
160	8192	57.0	10	8.6s	-253
160	16384	60.7	10	35.4s	-270

Approximate square roots of $\Sigma_2 = \sigma'^2 \cdot g^{-T} \cdot g - r^2$ on Intel Core i7-4850HQ CPU at 2.30GHz, 2 cores for Denman-Beavers, 4 cores for estimating the scaling factor, one core for sampling.

Thank You

Code <https://bitbucket.org/malb/gghlite-flint>

Paper <http://ia.cr/2014/928>