So how hard is solving hard lattice problems anyway?

Martin R. Albrecht @martinralbrecht

10 July 2019, AfricaCrypt

Based on joint work with Alex Davidson, Amit Deo, Benjamin R. Curtis, Eamonn W. Postlethwaite, Elena Kirshanova, Fernando Virdia, Florian Göpfert, Gottfried Herold, John M. Schanck, Léo Ducas, Marc Stevens, Rachel Player, Sam Scott, Thomas Wunderer and Vlad Gheorghiu as well as the work of many other authors.

Introduction

NIST Round 1: Selected Cost Estimates

Cost Model \ Scheme	Kyber	NewHope	NTRU HRSS	SNTRU'
$0.292\,\beta$[1]	180	259	136	155
$1/(2e)\,\beta\log(\beta) - \beta + 16.1$[2]	456	738	313	370
$1/8\,\beta\log(\beta) - 0.75\beta + 2.3$[3]	248	416	165	200
$0.265\,\beta$[1]	163	235	123	140
$1/(4e)\,\beta\log(\beta) - 1/2\beta + 8$	228	369	157	187

Source: Martin R. Albrecht, Benjamin R. Curtis, Amit Deo, Alex Davidson, Rachel Player, Eamonn W. Postlethwaite, Fernando Virdia, and Thomas Wunderer. Estimate All the LWE, NTRU Schemes! In: *SCN 18*. Ed. by Dario Catalano and Roberto De Prisco. Vol. 11035. LNCS. Springer, Heidelberg, Sept. 2018, pp. 351–367. DOI: 10.1007/978-3-319-98113-0_19, https://estimate-all-the-lwe-ntru-schemes.github.io/docs/

[1] Erdem Alkim, Léo Ducas, Thomas Pöppelmann, and Peter Schwabe. Post-quantum Key Exchange - A New Hope. In: *USENIX Security 2016*. Ed. by Thorsten Holz and Stefan Savage. USENIX Association, Aug. 2016, pp. 327–343

[2] Martin R. Albrecht, Rachel Player, and Sam Scott. On the concrete hardness of Learning with Errors. In: *Journal of Mathematical Cryptology* 9.3 (2015), pp. 169–203

[3] Le Trieu Phong, Takuya Hayashi, Yoshinori Aono, and Shiho Moriai. LOTUS. Tech. rep. available at https://csrc.nist.gov/projects/post-quantum-cryptography/round-1-submissions. National Institute of Standards and Technology, 2017

Given $(\mathbf{A}, \mathbf{c})$, find $\mathbf{s}$ when

$$\begin{pmatrix} \\ \mathbf{c} \\ \\ \end{pmatrix} \equiv \begin{pmatrix} \leftarrow & n & \rightarrow \\ & & \\ & \mathbf{A} & \\ & & \end{pmatrix} \cdot \begin{pmatrix} \\ \mathbf{s} \\ \\ \end{pmatrix} + \begin{pmatrix} \\ \mathbf{e} \\ \\ \end{pmatrix}$$

for $\mathbf{c} \in \mathbb{Z}_q^m$, $\mathbf{A} \in \mathbb{Z}_q^{m \times n}$, and $\mathbf{s} \in \mathbb{Z}^n$ and $\mathbf{e} \in \mathbb{Z}^m$ having small coefficients.

(Matrix-)NTRU

Let $\mathbf{F}, \mathbf{G}$ be two $n \times n$ matrices over $\mathbb{Z}_q$ with short entries. Given

$$\mathbf{H} \equiv \mathbf{F}^{-1} \cdot \mathbf{G}$$

find (a small multiple of) $\mathbf{F}$ or $\mathbf{G}$.

(Matrix-)NTRU

Let $\mathbf{F}, \mathbf{G}$ be two $n \times n$ matrices over $\mathbb{Z}_q$ with short entries. Given

$$\mathbf{H} \equiv \mathbf{F}^{-1} \cdot \mathbf{G}$$

find (a small multiple of) $\mathbf{F}$ or $\mathbf{G}$.

Note

I will focus on LWE in this talk, but the techniques translate (with some modifications) to NTRU.

Primal Approach

Unique SVP Approach

We can reformulate $\mathbf{c} - \mathbf{A} \cdot \mathbf{s} \equiv \mathbf{e} \bmod q$ over the Integers as:

$$\begin{pmatrix} q\mathbf{I} & -\mathbf{A} \\ 0 & \mathbf{I} \end{pmatrix} \cdot \begin{pmatrix} * \\ \mathbf{s} \end{pmatrix} + \begin{pmatrix} \mathbf{c} \\ \mathbf{0} \end{pmatrix} = \begin{pmatrix} \mathbf{e} \\ \mathbf{s} \end{pmatrix}$$

Alternatively:

$$\mathbf{B} = \begin{pmatrix} q\mathbf{I} & -\mathbf{A} & \mathbf{c} \\ 0 & \mathbf{I} & 0 \\ 0 & 0 & 1 \end{pmatrix}, \qquad \mathbf{B} \cdot \begin{pmatrix} * \\ \mathbf{s} \\ 1 \end{pmatrix} = \begin{pmatrix} \mathbf{e} \\ \mathbf{s} \\ 1 \end{pmatrix}$$

In other words, there exists an integer-linear combination of the columns of $\mathbf{B}$ that produces a vector with “unusually” small coefficients $\rightarrow$ a unique shortest vector.

Unique Shortest Vector Problem

Find a unique shortest vector amongst the integer combinations of the columns of:

$$\mathbf{B} = \begin{pmatrix} q\mathbf{I} & -\mathbf{A} & \mathbf{c} \\ 0 & \mathbf{I} & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

where $\mathbf{B} \in \mathbb{Z}^{d \times d}$.

Unique Shortest Vector Problem

Find a unique shortest vector amongst the integer combinations of the columns of:

$$\mathbf{B} = \begin{pmatrix} q\mathbf{I} & -\mathbf{A} & \mathbf{c} \\ 0 & \mathbf{I} & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

where $\mathbf{B} \in \mathbb{Z}^{d \times d}$.

NTRU

In the case of LWE we have (up to $\pm$) one such short vector. In the case of NTRU we have n.

Lattice Reduction

Lattices

A lattice is a discrete subgroup of $\mathbb{R}^d$ and can be written as $\Lambda = \{\sum_{i=1}^{d} v_i \cdot \mathbf{b}_i \mid v_i \in \mathbb{Z}\}$

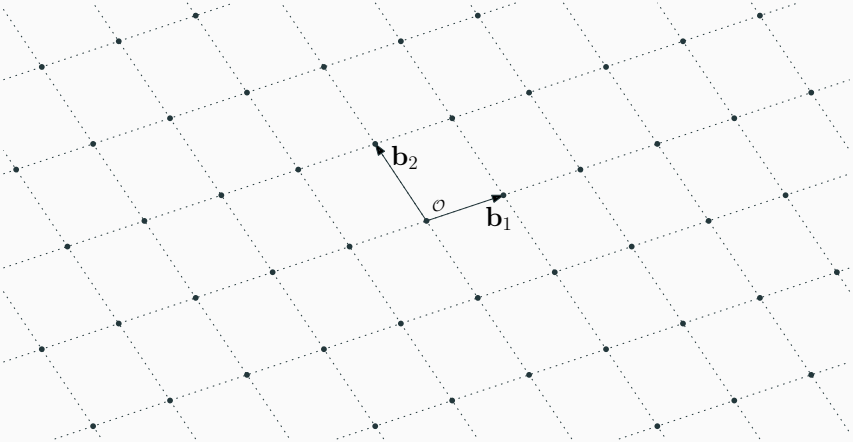

Picture Credit: Joop van der Pol

A Tale of Bad Bases . . .

With “bad basis” finding short vectors assumed hard.

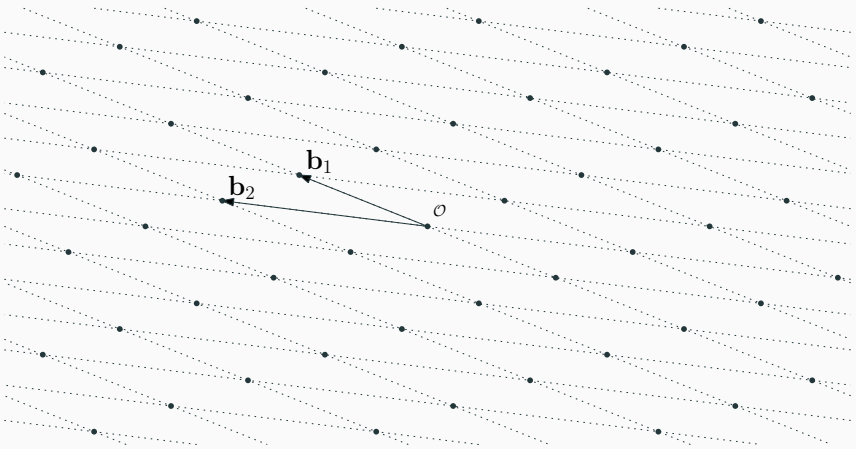

Picture Credit: Joop van der Pol

. . . and Good Bases

With a “good basis” many lattice problems are easy.

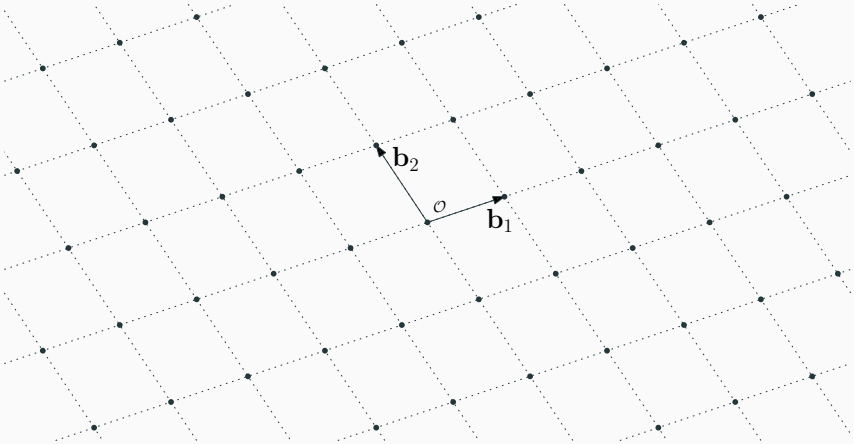

Picture Credit: Joop van der Pol

Lattice Volume

The volume of a lattice is the volume of its fundamental parallelepiped.

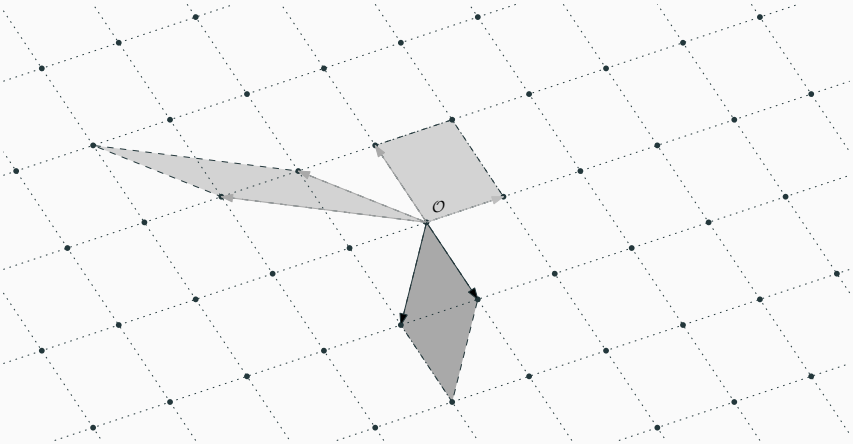

Picture Credit: Joop van der Pol

Gaussian Heuristic

The shortest vector in the lattice has expected norm

$$\lambda_1(\Lambda) \approx \sqrt{\frac{d}{2\pi e}}\mathrm{Vol}(\Lambda)^{1/d}.$$

Unusually Shortest Vector

When $\lambda_1(\Lambda) \ll \sqrt{\frac{d}{2\pi e}}\mathrm{Vol}(\Lambda)^{1/d}$.

Length of Gram-Schmidt Vectors

It will be useful to consider the lengths of the Gram-Schmidt vectors.

The vector $\mathbf{b}_i^*$ is the orthogonal projection of $\mathbf{b}_i$ to the space spanned by the vectors $\mathbf{b}_0, \ldots, \mathbf{b}_{i-1}$.

Informally, this means taking out the contributions in the directions of previous vectors $\mathbf{b}_0, \ldots, \mathbf{b}_{i-1}$.

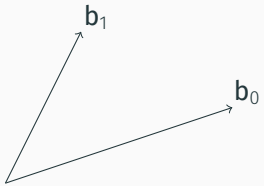

Length of Gram-Schmidt Vectors

It will be useful to consider the lengths of the Gram-Schmidt vectors.

The vector $\mathbf{b}_i^*$ is the orthogonal projection of $\mathbf{b}_i$ to the space spanned by the vectors $\mathbf{b}_0, \ldots, \mathbf{b}_{i-1}$.

Informally, this means taking out the contributions in the directions of previous vectors $\mathbf{b}_0, \ldots, \mathbf{b}_{i-1}$.

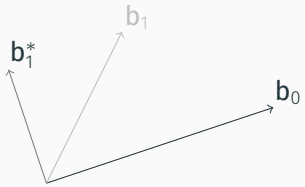

Example

```
sage: A = IntegerMatrix.random(120, "qary", k=60, bits=20)[::-1]
sage: M = GSO.Mat(A); M.update_gso()
sage: lg = [(i,log(r_, 2)/2) for i, r_ in enumerate(M.r())]
sage: line(lg, **plot_kwds)
```

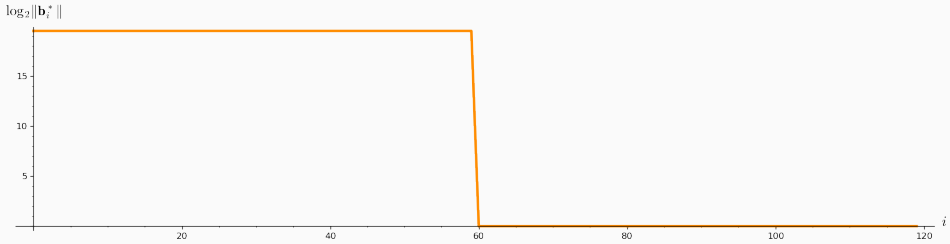

Example - LLL

```
sage: A = LLL.reduction(A)
sage: M = GSO.Mat(A); M.update_gso()
sage: lg = [(i,log(r_, 2)/2) for i, r_ in enumerate(M.r())]
sage: line(lg, **plot_kwds)
```

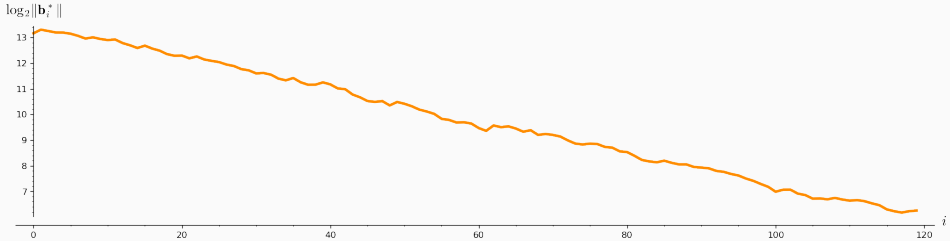

Geometric Series Assumption: The shape after lattice reduction is a line with a flatter slope as lattice reduction gets stronger.

Strong Lattice Reduction: BKZ Algorithm

$$\overbrace{\left(\begin{array}{cccccccccc} & & & & & & & & & \\ & & & & & & & & & \\ \mathbf{b}_0 & \mathbf{b}_1 & \mathbf{b}_2 & \mathbf{b}_3 & \mathbf{b}_4 & \mathbf{b}_5 & \mathbf{b}_6 & \mathbf{b}_7 & \dots \\ & & & & & & & & & \\ & & & & & & & & & \end{array}\right)}^{\beta = 5}$$

Picture credit: Eamonn Postlethwaite

Strong Lattice Reduction: BKZ Algorithm

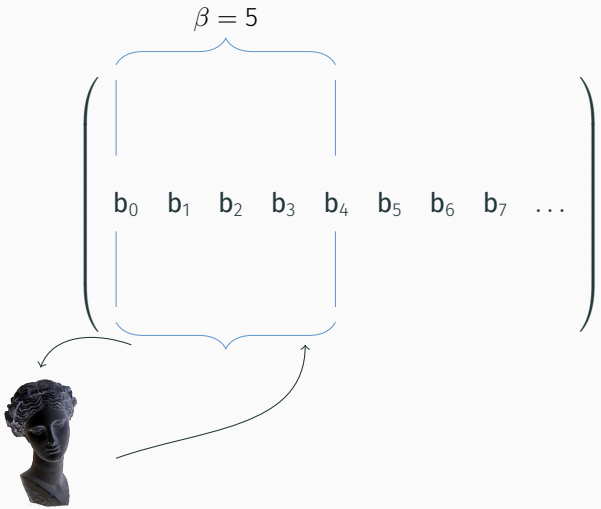

Picture credit: Eamonn Postlethwaite

Strong Lattice Reduction: BKZ Algorithm

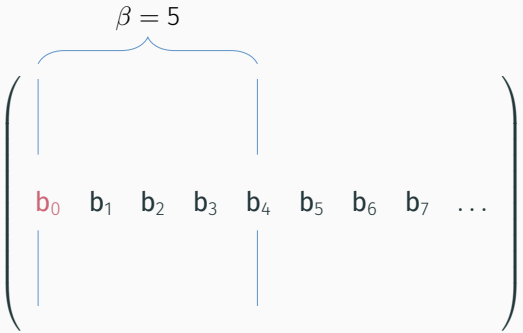

Picture credit: Eamonn Postlethwaite

Strong Lattice Reduction: BKZ Algorithm

$$\beta = 5$$

$$\begin{pmatrix} & & & & & & & & \\ \mathbf{b}_0 & \mathbf{b}_1 & \mathbf{b}_2 & \mathbf{b}_3 & \mathbf{b}_4 & \mathbf{b}_5 & \mathbf{b}_6 & \mathbf{b}_7 & \dots \\ & & & & & & & & \end{pmatrix}$$

Picture credit: Eamonn Postlethwaite

Strong Lattice Reduction: BKZ Algorithm

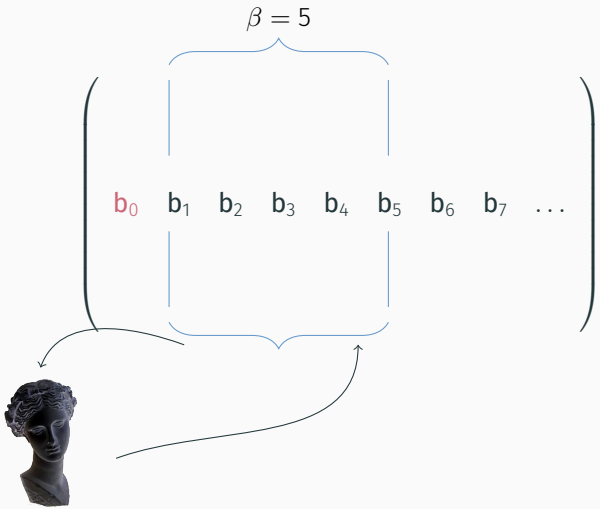

Picture credit: Eamonn Postlethwaite

Strong Lattice Reduction: BKZ Algorithm

$$\beta = 5$$

$$\begin{pmatrix} | & | & | & | & | & | & | & | & \\ \mathbf{b}_0 & \mathbf{b}_1 & \mathbf{b}_2 & \mathbf{b}_3 & \mathbf{b}_4 & \mathbf{b}_5 & \mathbf{b}_6 & \mathbf{b}_7 & \dots \\ | & | & | & | & | & | & | & | & \end{pmatrix}$$

Picture credit: Eamonn Postlethwaite

Strong Lattice Reduction: BKZ Algorithm

$$\beta = 5$$

$$\begin{pmatrix} & & & & & & & & \\ \mathbf{b}_0 & \mathbf{b}_1 & \mathbf{b}_2 & \mathbf{b}_3 & \mathbf{b}_4 & \mathbf{b}_5 & \mathbf{b}_6 & \mathbf{b}_7 & \dots \\ & & & & & & & & \end{pmatrix}$$

Picture credit: Eamonn Postlethwaite

Strong Lattice Reduction: BKZ Algorithm

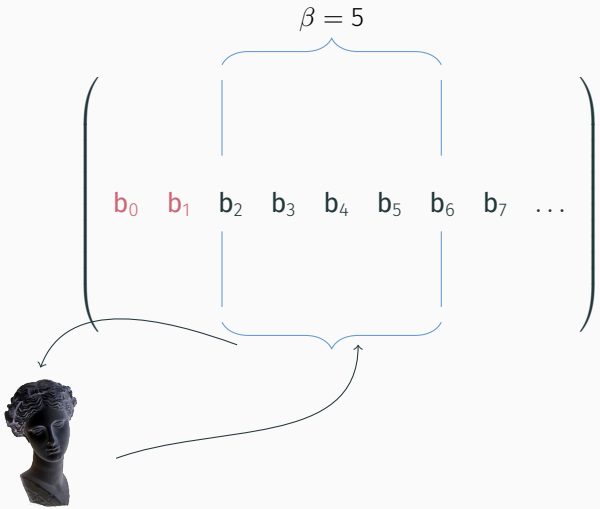

Picture credit: Eamonn Postlethwaite

Strong Lattice Reduction: BKZ Algorithm

$$\beta = 5$$

$$\left(\begin{array}{ccccccccc} \mathbf{b}_0 & \mathbf{b}_1 & \mathbf{b}_2 & \mathbf{b}_3 & \mathbf{b}_4 & \mathbf{b}_5 & \mathbf{b}_6 & \mathbf{b}_7 & \dots \end{array} \right)$$

Picture credit: Eamonn Postlethwaite

BKZ Algorithm

Data: LLL-reduced lattice basis $\mathbf{B}$
Data: block size β
repeat *until no more change*
  for $\kappa \leftarrow 0$ **to** $d - 1$ **do**
    LLL on local projected block $[\kappa, \ldots, \kappa + \beta - 1]$;
    $\mathbf{v} \leftarrow$ find shortest vector in local projected block $[\kappa, \ldots, \kappa + \beta - 1]$;
    insert $\mathbf{v}$ into $\mathbf{B}$;
  end

Jargon

An outer loop iteration is called a “tour”.

Slope

The slope depends on the **root Hermite factor** δ which depends on the block size β.

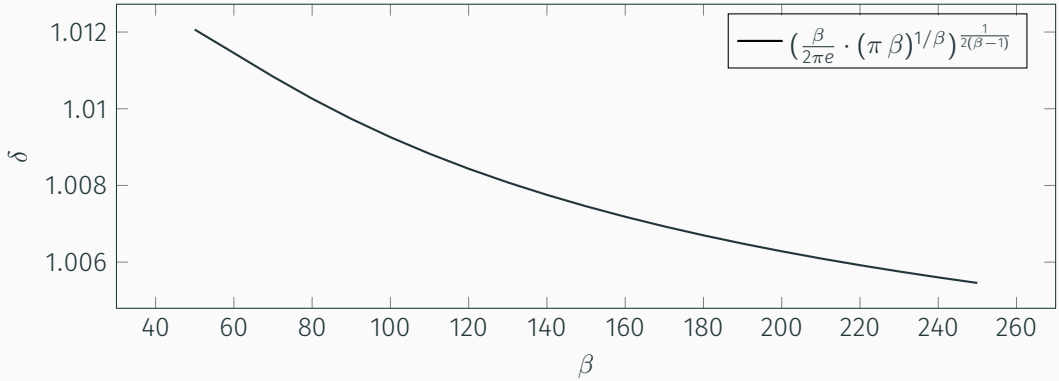

Yuanmi Chen. Réduction de réseau et sécurité concrète du chiffrement complètement homomorphe. PhD thesis. Paris 7, 2013

Behaviour in Practice: BKZ-60 in Dimension 180 I

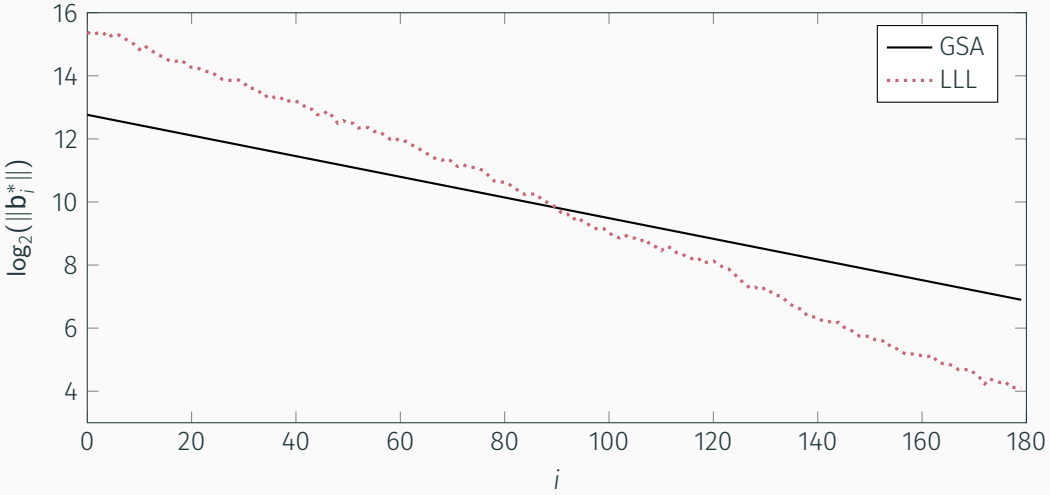

Behaviour in Practice: BKZ-60 in Dimension 180 II

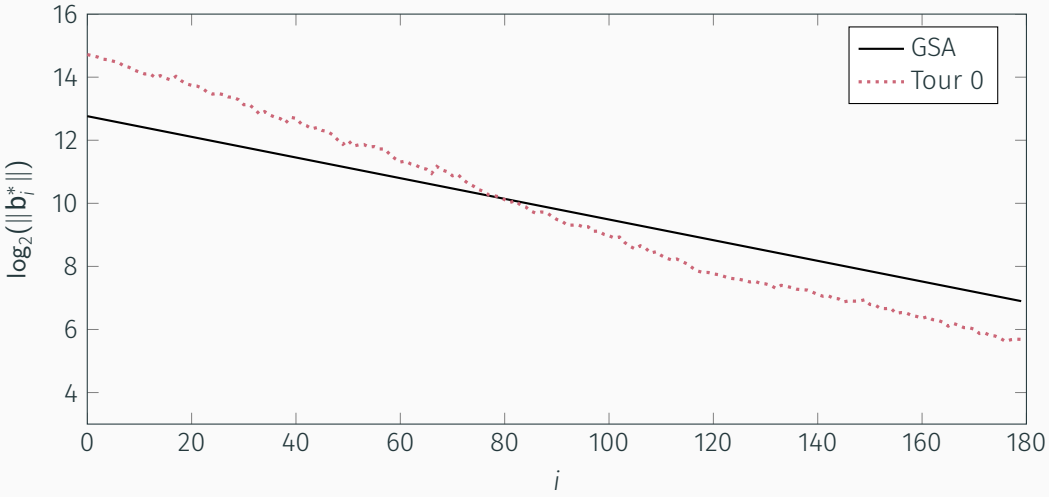

Behaviour in Practice: BKZ-60 in Dimension 180 III

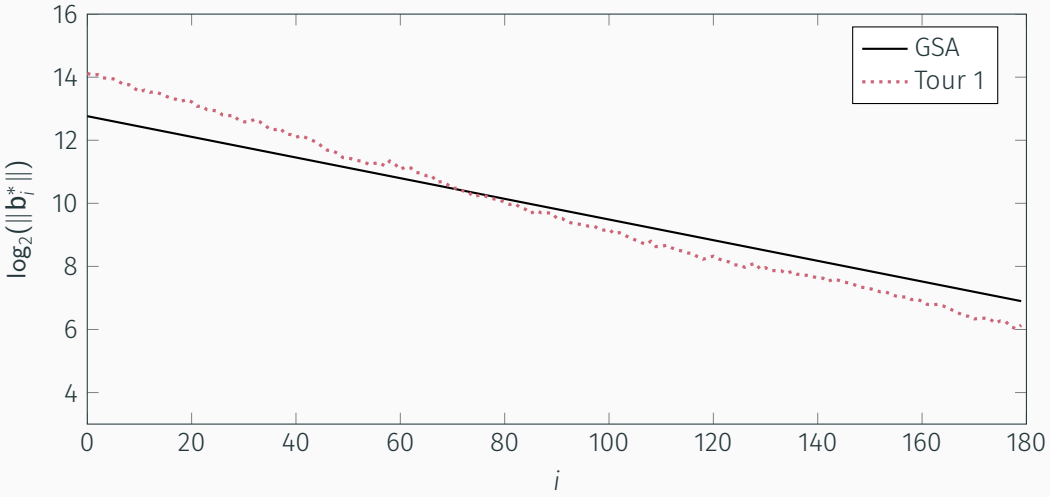

Behaviour in Practice: BKZ-60 in Dimension 180 iv

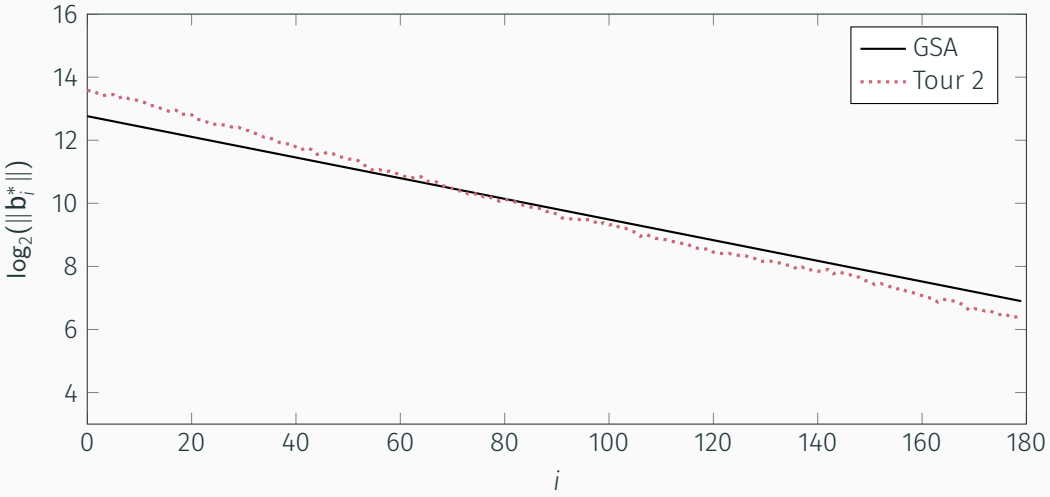

Behaviour in Practice: BKZ-60 in Dimension 180 v

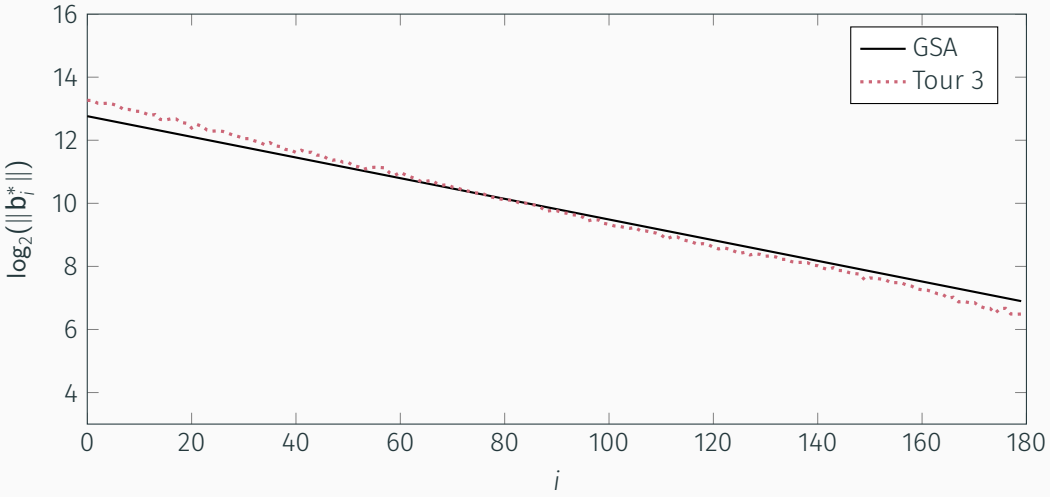

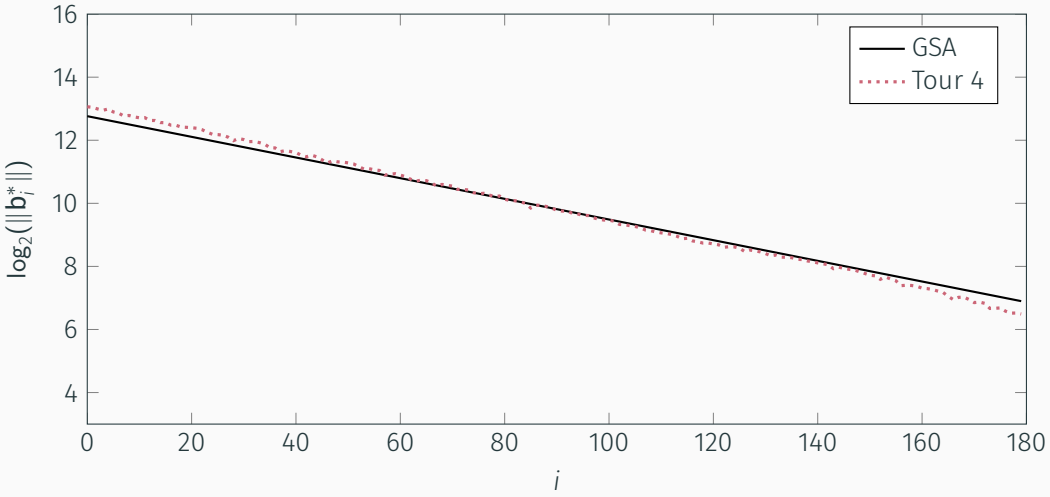
GSA
Tour 4
$\log_2(\|\mathbf{b}_i^*\|)$
i
16
14
12
10
8
6
4
0
20
40
60
80
100
120
140
160
180

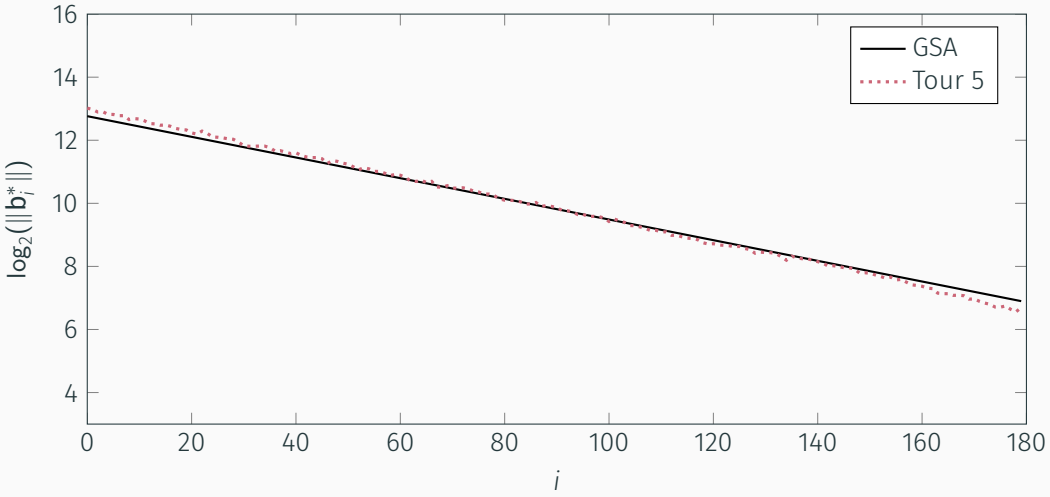
GSA
Tour 5
$\log_2(\|\mathbf{b}_i^*\|)$
i
16
14
12
10
8
6
4
0
20
40
60
80
100
120
140
160
180

Behaviour in Practice: BKZ-60 in Dimension 180 viii

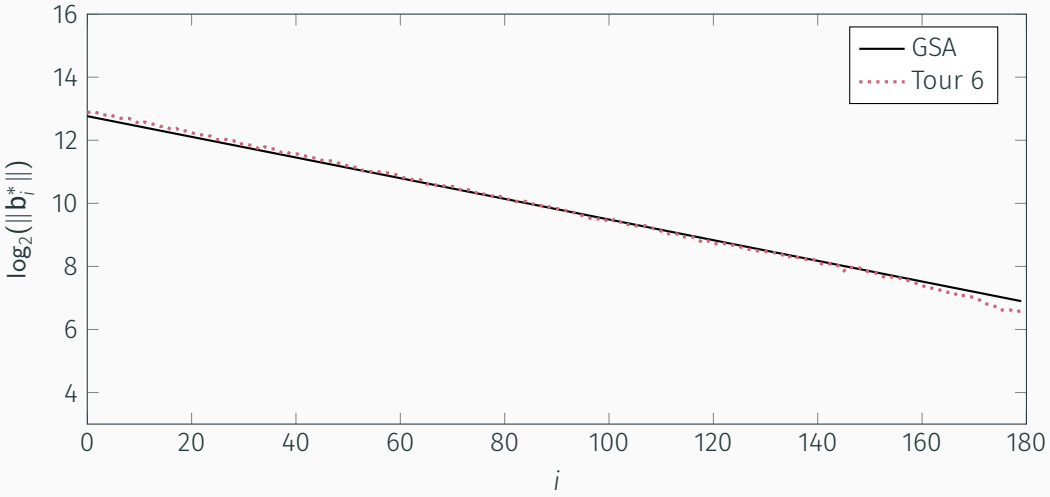

Behaviour in Practice: BKZ-60 in Dimension 180 ix

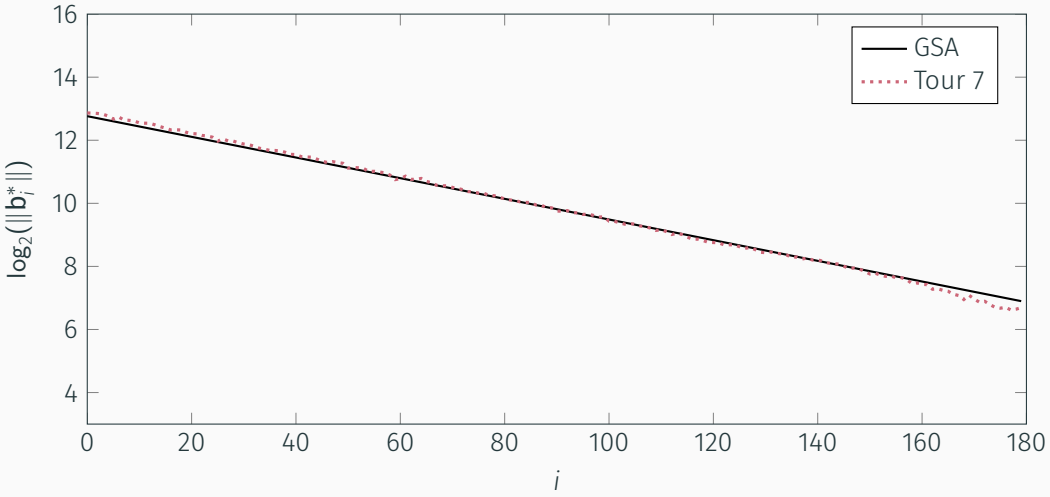

Behaviour in Practice: BKZ-60 in Dimension 180 x

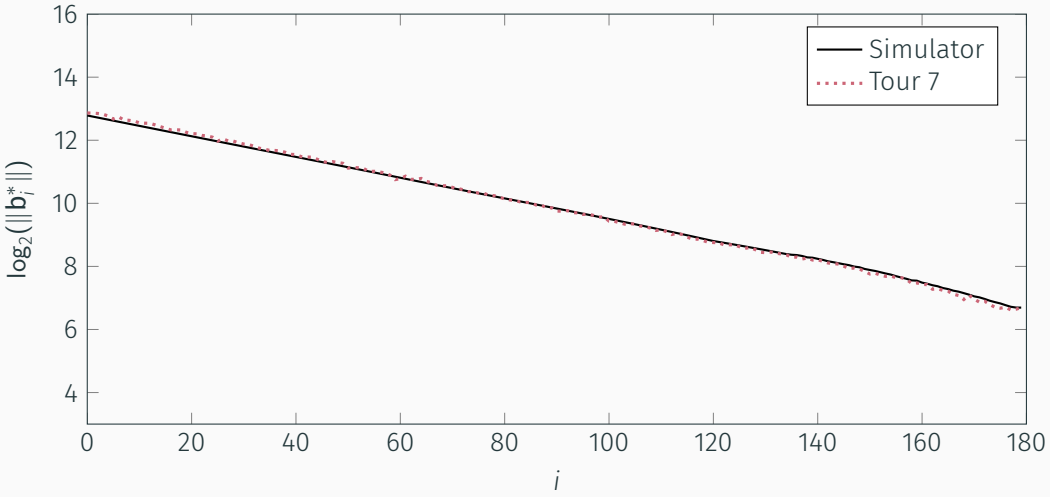

Success Condition for uSVP (Expectation)

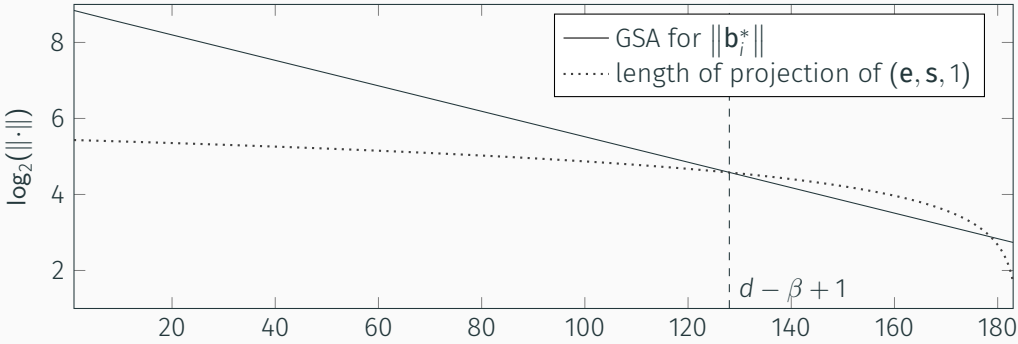

Erdem Alkim, Léo Ducas, Thomas Pöppelmann, and Peter Schwabe. Post-quantum Key Exchange - A New Hope. In: *USENIX Security 2016*. Ed. by Thorsten Holz and Stefan Savage. USENIX Association, Aug. 2016, pp. 327–343

Success Condition for uSVP (Observed)

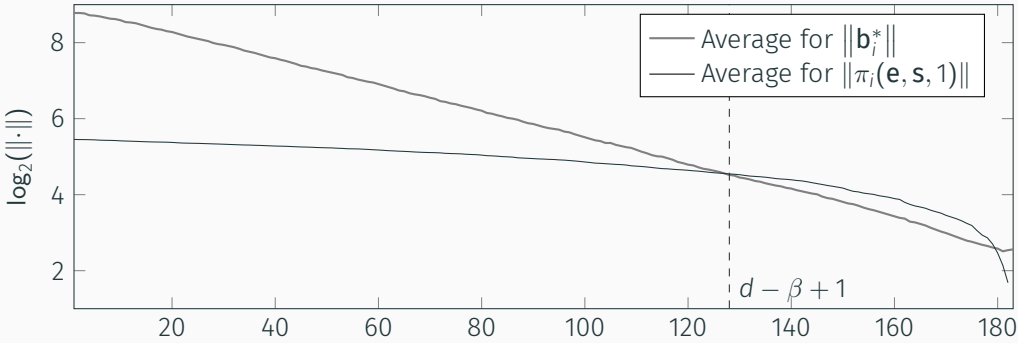

Martin R. Albrecht, Florian Göpfert, Fernando Virdia, and Thomas Wunderer. Revisiting the Expected Cost of Solving uSVP and Applications to LWE. In: *ASIACRYPT 2017, Part I*. ed. by Tsuyoshi Takagi and Thomas Peyrin. Vol. 10624. LNCS. Springer, Heidelberg, Dec. 2017, pp. 297–322. DOI: 10.1007/978-3-319-70694-8_11

Solving SVP

Solving SVP

Cost Model \ Scheme	Kyber	NewHope	NTRU HRSS	SNTRU'
$0.292\,\beta$[1]	180	259	136	155
$1/(2e)\,\beta\log(\beta) - \beta + 16.1$[2]	456	738	313	370
$1/8\,\beta\log(\beta) - 0.75\beta + 2.3$[3]	248	416	165	200
$0.265\,\beta$[1]	163	235	123	140
$1/(4e)\,\beta\log(\beta) - 1/2\beta + 8$	228	369	157	187

Sieving

- Produce new, shorter vectors by considering sums and differences of existing vectors
- **Time:** $2^{\Theta(\beta)}$
- **Memory:** $2^{\Theta(\beta)}$

Enumeration

- Search through vectors smaller than a given bound: project down to 1-dim problem, lift to 2-dim problem ...
- **Time:** $2^{\Theta(\beta\log\beta)}$
- **Memory:** $\mathrm{poly}(\beta)$

Enumeration Estimates

The first estimate extrapolates a dataset from [Che13]

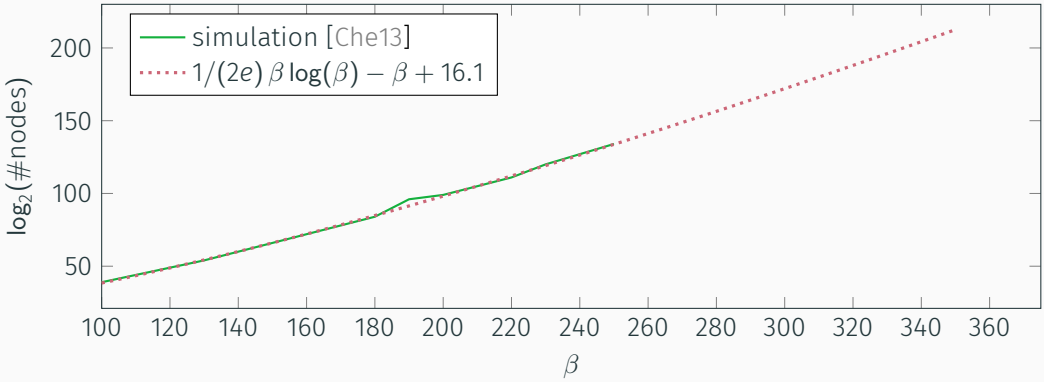

Extended Enumeration Simulation

That estimate compared to our simulation

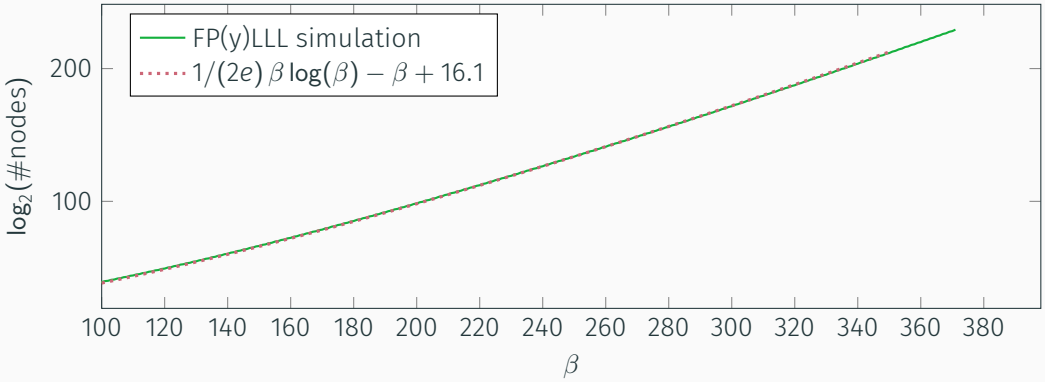

Enumeration Simulation vs Experiments

Assuming 1 node $\approx$ 100 cpu cycles:

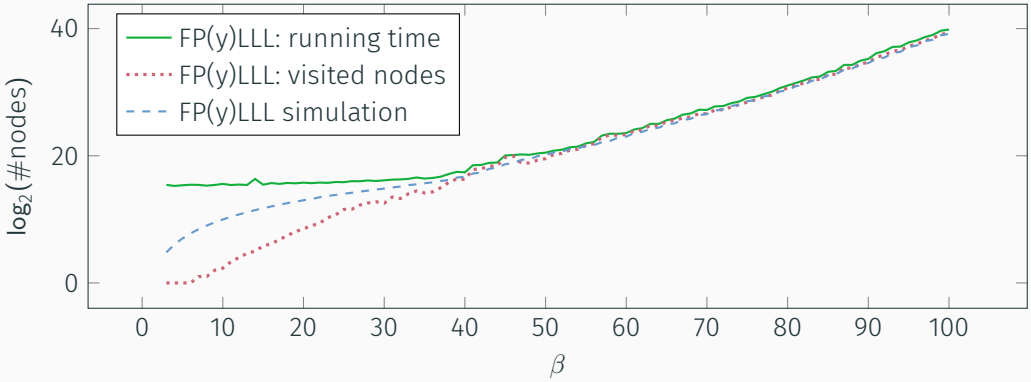

Enumeration Worst-Case Complexity

Cost Model \ Scheme	Kyber	NewHope	NTRU HRSS	SNTRU'
$1/(2e)\,\beta\log(\beta) - \beta + 16.1$[2]	456	738	313	370
$1/8\,\beta\log(\beta) - 0.75\beta + 2.3$[3]	248	416	165	200

"We obtain a new worst-case complexity upper bound, as well as the first worst-case complexity lower bound, both of the order d of $2^{O(d)} \cdot d^{\frac{d}{2e}}$ (up to polynomial factors) bit operations, where d is the rank of the lattice."[4]

[4] Guillaume Hanrot and Damien Stehlé. Improved Analysis of Kannan's Shortest Lattice Vector Algorithm. In: *CRYPTO 2007*. Ed. by Alfred Menezes. Vol. 4622. LNCS. Springer, Heidelberg, Aug. 2007, pp. 170–186. DOI: 10.1007/978-3-540-74143-5_10.

Enumeration Heuristic Best-Case Complexity

Cost Model \ Scheme	Kyber	NewHope	NTRU HRSS	SNTRU'
$1/(2e)\,\beta\log(\beta) - \beta + 16.1$[2]	456	738	313	370
$1/8\,\beta\log(\beta) - 0.75\beta + 2.3$[3]	248	416	165	200

"This suggests that, independently of the quality of the reduced basis, the complexity of enumeration will be at least $d^{d/8}$ polynomial-time operations for many lattices."[5]

[5] Phong Q. Nguyen. Hermite's Constant and Lattice Algorithms. In: ed. by Phong Q. Nguyen and Brigitte Vallée. ISC. Springer, Heidelberg, 2010, pp. 19–69. ISBN: 978-3-642-02294-4. DOI: 10.1007/978-3-642-02295-1.

$1/8 \approx 0.125$ **v** $1/(2e) \approx 0.184$

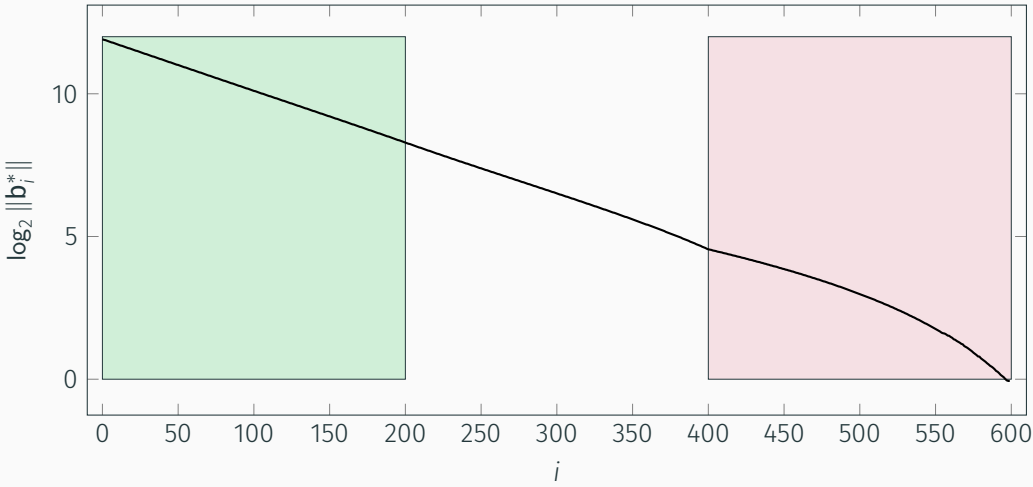

Why we can't have Nice Things

- We run enumeration many times each succeeding with low probability of success and re-randomise in between: this destroys the nice GSA-line shape
- Thus, before enumerating a local block, we run some local preprocessing with some block size $\beta' < \beta$
- In the sandpile model,[6] as the algorithm proceeds through the indices i, a "bump" accumulates from index $i + 1$ onward.

[6] Guillaume Hanrot, Xavier Pujol, and Damien Stehlé. Analyzing Blockwise Lattice Algorithms Using Dynamical Systems. In: *CRYPTO 2011*. Ed. by Phillip Rogaway. Vol. 6841. LNCS. Springer, Heidelberg, Aug. 2011, pp. 447–464. DOI: 10.1007/978-3-642-22792-9_25.

Idea: Overshoot Preprocessing (WIP)

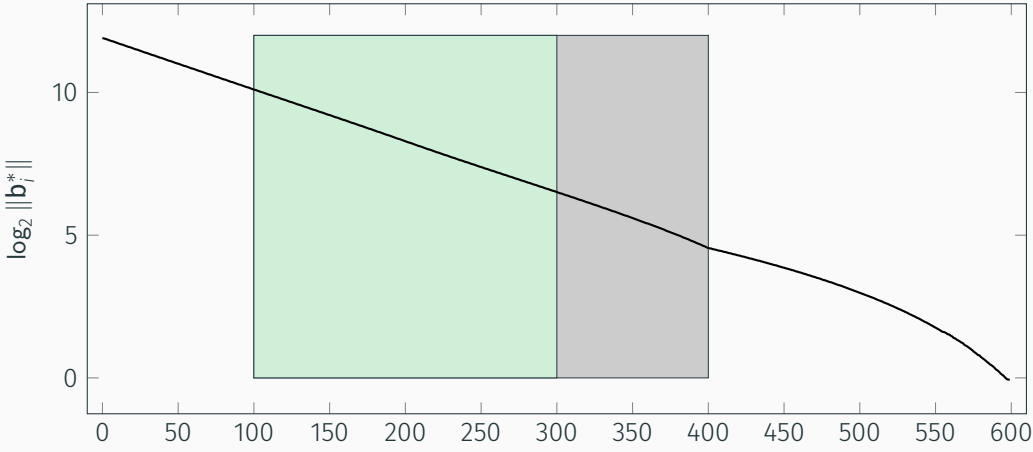

Preprocessing in dimension $(1 + c) \cdot \beta$ for enumeration in dimension β.[7]

[7] Joint work with Shi Bai, Léo Ducas and Damien Stehlé

The Tail Problem

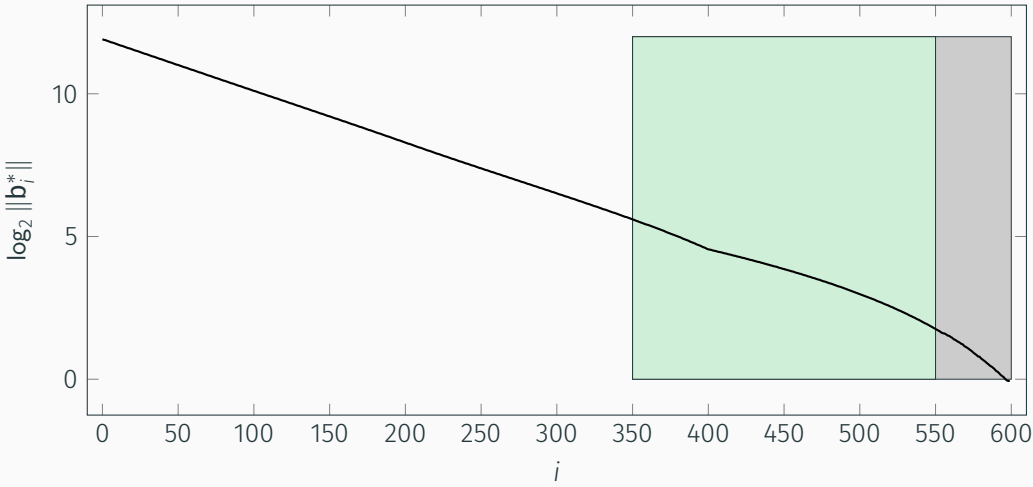

Handling the Tail: Reduce Block Size

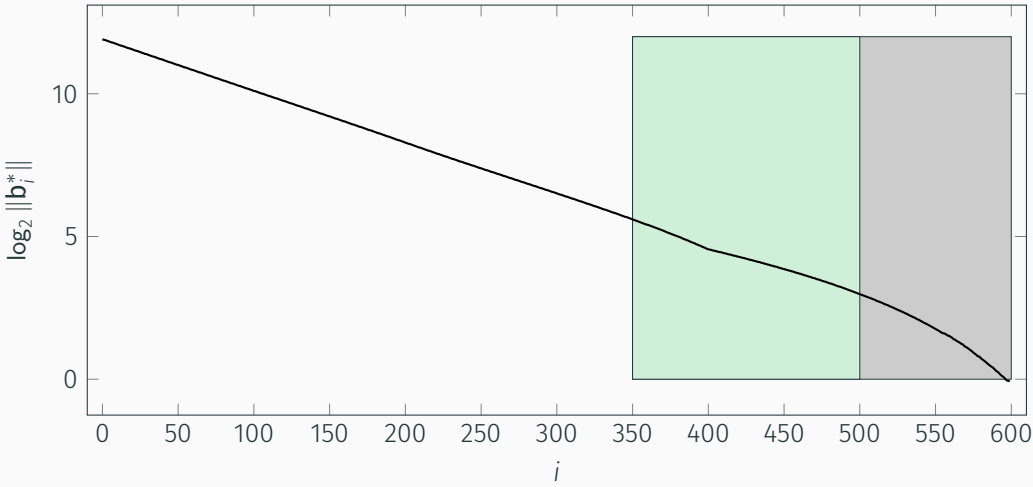

Quality Degradation ($d = 2\beta$) (WIP)

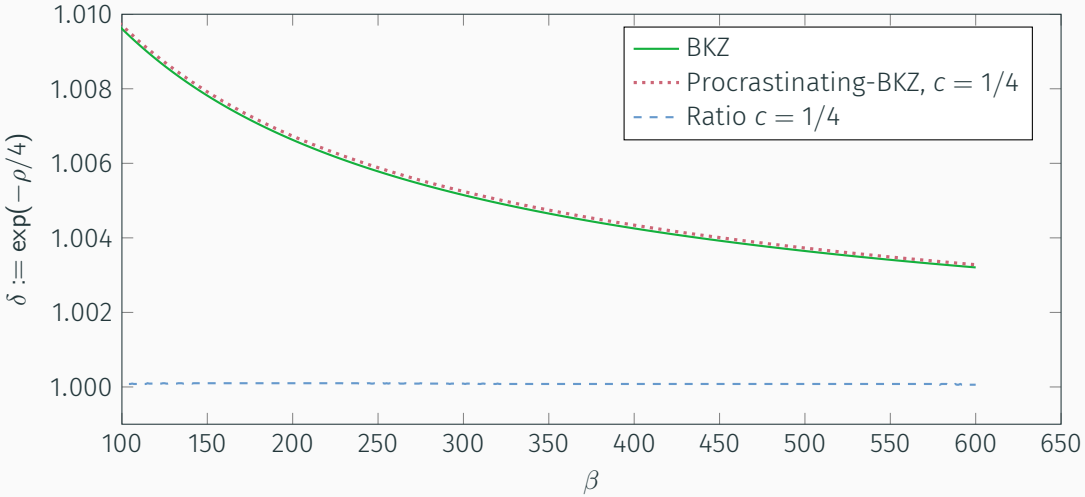

Idea: Self-Dual Procrastinating-BKZ (WIP)

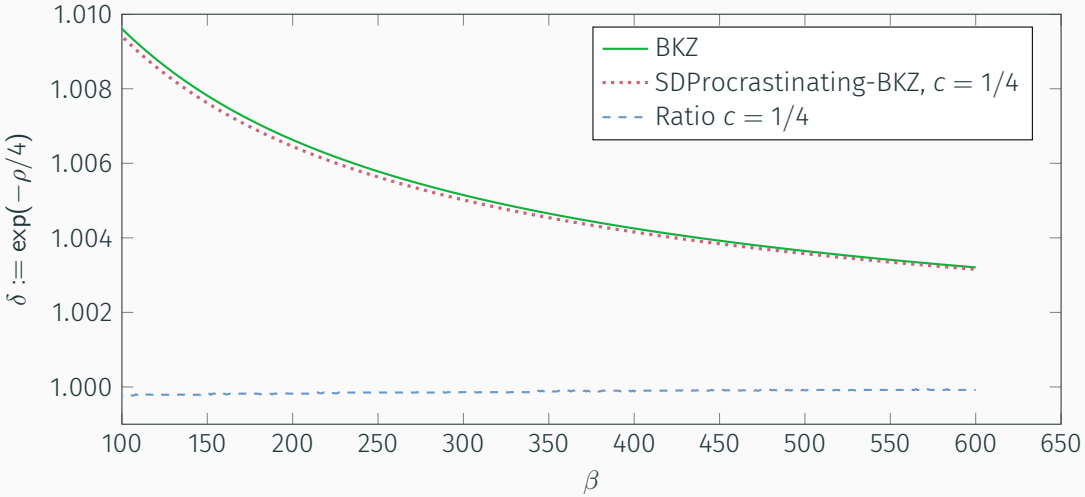

Performance (WIP)

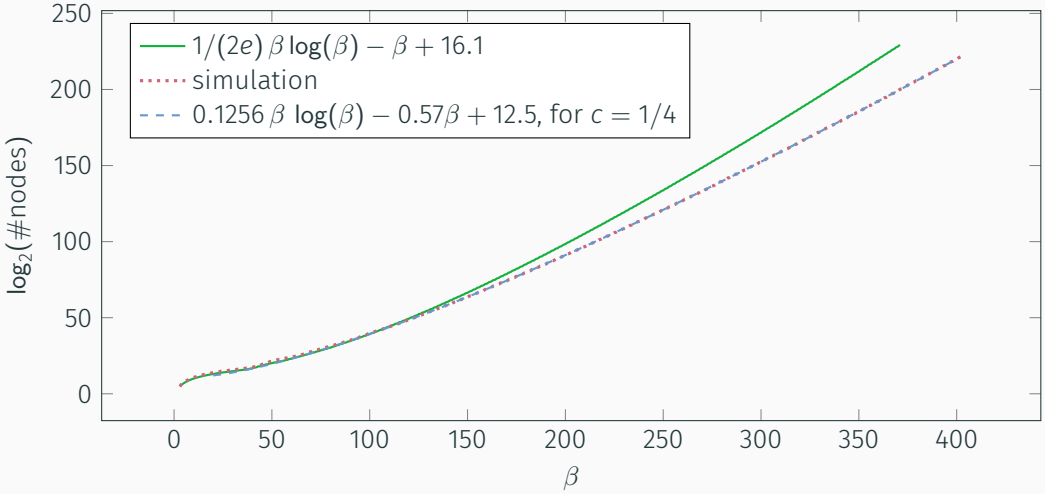

Sieving vs Enumeration

Cost Model \ Scheme	Kyber	NewHope	NTRU HRSS	SNTRU'
$0.292\,\beta$[1]	180	259	136	155
$1/(2e)\,\beta\log(\beta) - \beta + 16.1$[2]	456	738	313	370
$1/8\,\beta\log(\beta) - 0.75\beta + 2.3$[3]	248	416	165	200

Sieving is asymptotically faster than enumeration, but does it beat enumeration in practical or cryptographic dimensions?

Sieving: Key Ideas I

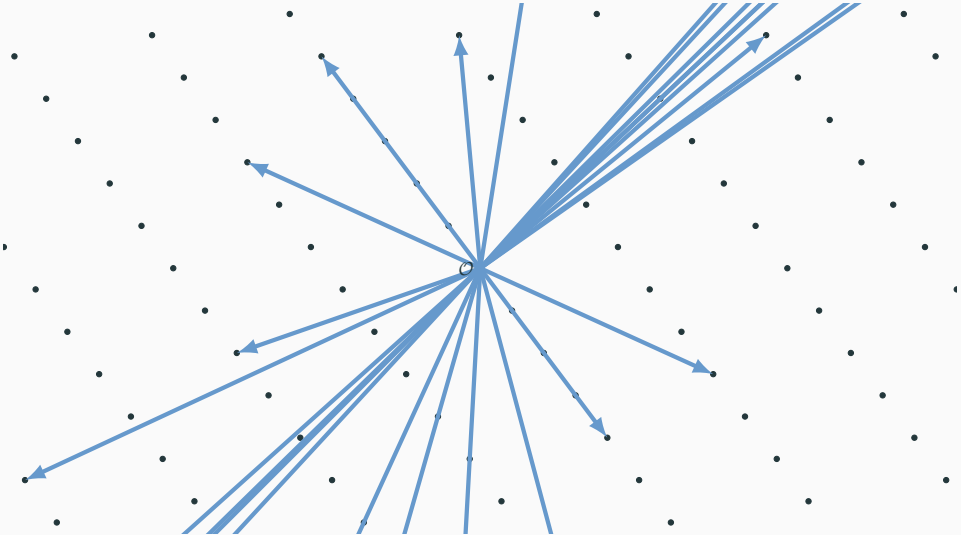

Sieving: Key Ideas II

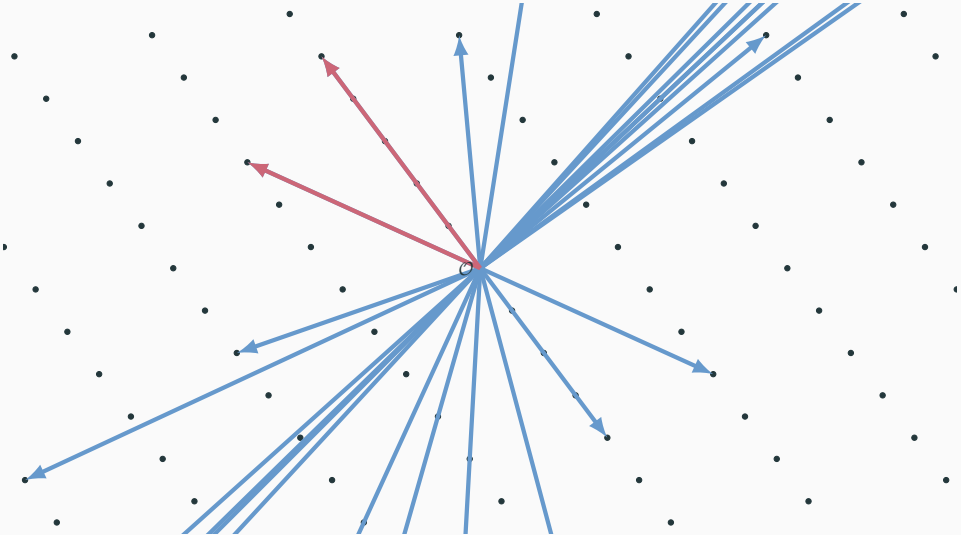

Sieving: Key Ideas III

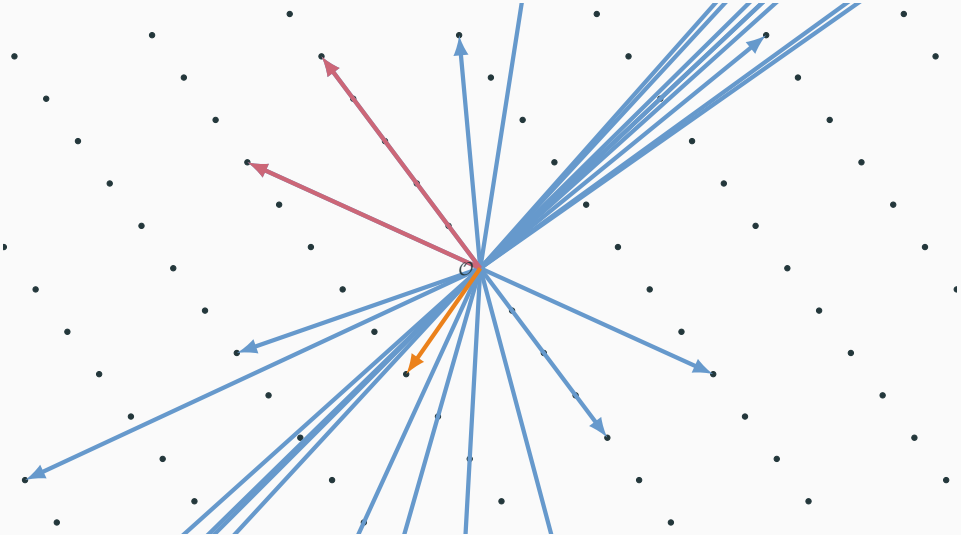

Sieving: Popcount I

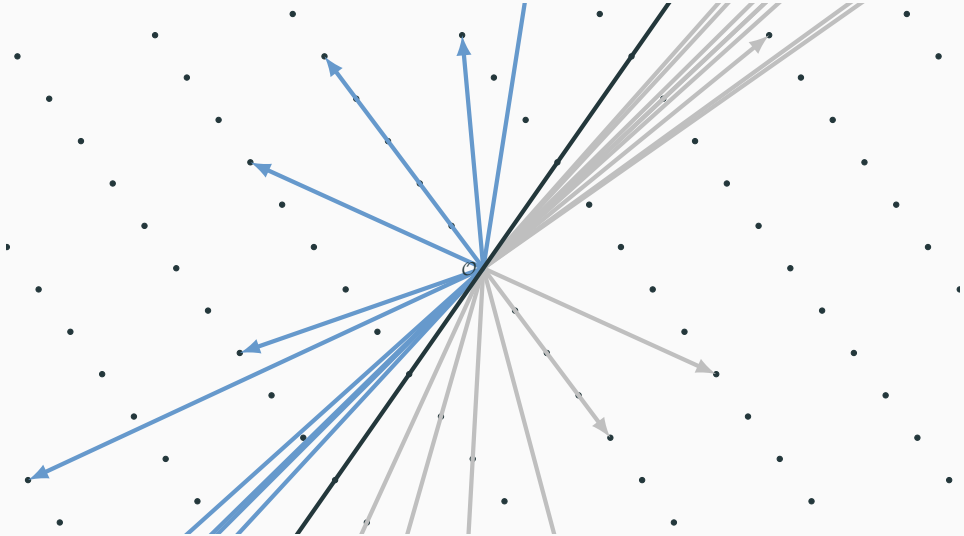

Sieving: Popcount II

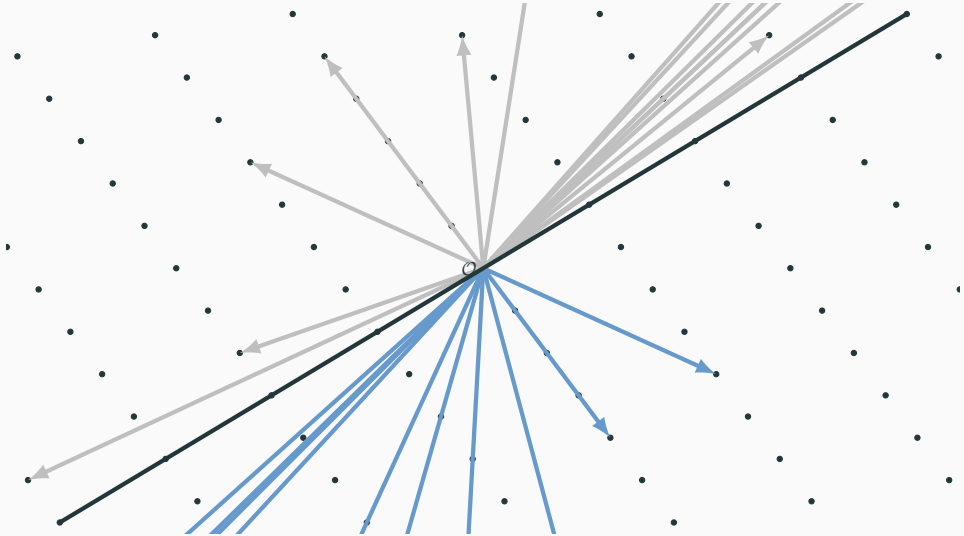

Sieving: Popcount III

- For a given plane, denote a vector being on the “left” as 0, being on the “right” as 1.
- This defines a 1-bit locality sensitive hash (LSH) function.
- Consider many such hash functions and concatenate their output.
- Two vectors are close if they agree on many bits of their hashes
- Comparison operation: XOR hash values and compute Hamming weight (“popcount”).

Sieving: Buckets

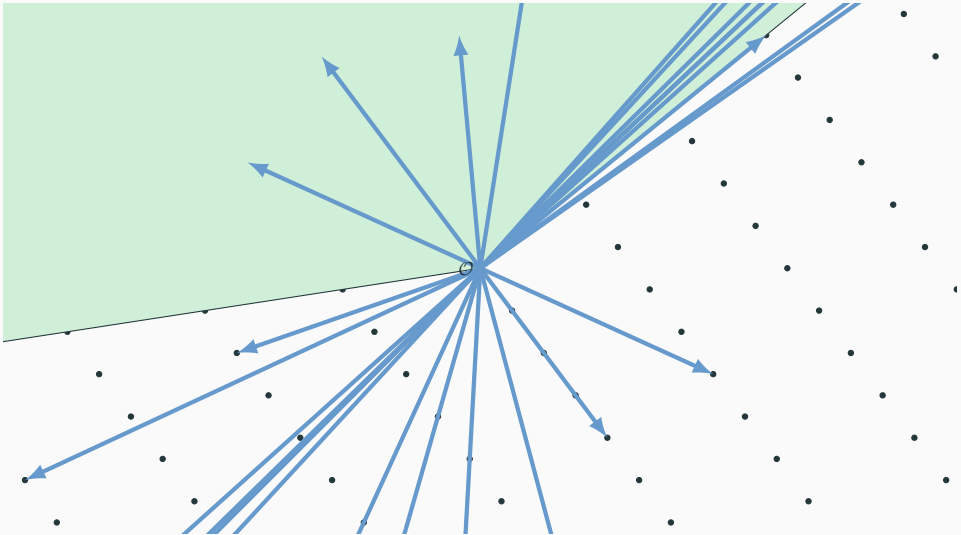

Sieving: Some Algorithms

Gauss Sample $(4/3)^{\beta/2+o(\beta)}$ vectors, compare them pairwise if they reduce to something shorter. **Cost**: $(4/3)^{\beta+o(\beta)} \approx 2^{0.41\,\beta+o(\beta)}$.[8]

BGJ Split search space into "buckets". **Cost**: $2^{0.311\,\beta+o(\beta)}$.[9]

BDGL Use codes to decide which bucket to consider. **Cost**: $2^{0.292\,\beta+o(\beta)}$.[10]

[8] Daniele Micciancio and Panagiotis Voulgaris. Faster Exponential Time Algorithms for the Shortest Vector Problem. In: *21st SODA*. ed. by Moses Charika. ACM-SIAM, Jan. 2010, pp. 1468–1480. DOI: 10.1137/1.9781611973075.119.

[9] Anja Becker, Nicolas Gama, and Antoine Joux. Speeding-up lattice sieving without increasing the memory, using sub-quadratic nearest neighbor search. Cryptology ePrint Archive, Report 2015/522. http://eprint.iacr.org/2015/522. 2015.

[10] Anja Becker, Léo Ducas, Nicolas Gama, and Thijs Laarhoven. New directions in nearest neighbor searching with applications to lattice sieving. In: *27th SODA*. ed. by Robert Krauthgamer. ACM-SIAM, Jan. 2016, pp. 10–24. DOI: 10.1137/1.9781611974331.ch2.

Sieving: G6K

G6K[11] is a Python/C++ framework for experimenting with sieving algorithms (inside and outside BKZ)

- Does not take the “oracle” view but considers sieves as stateful machines.
- Implements several sieve algorithms[12] (but not BDGL)
- Applies many recent tricks and adds new tricks for improving performance of sieving

[11] Martin R. Albrecht, Léo Ducas, Gottfried Herold, Elena Kirshanova, Eamonn W. Postlethwaite, and Marc Stevens. The General Sieve Kernel and New Records in Lattice Reduction. In: *EUROCRYPT 2019, Part II*. ed. by Yuval Ishai and Vincent Rijmen. Vol. 11477. LNCS. Springer, Heidelberg, May 2019, pp. 717–746. DOI: 10.1007/978-3-030-17656-3_25.

[12] Gauss, NV, BGJ1 (BGJ with one level of filtration)

Sieving: SVP

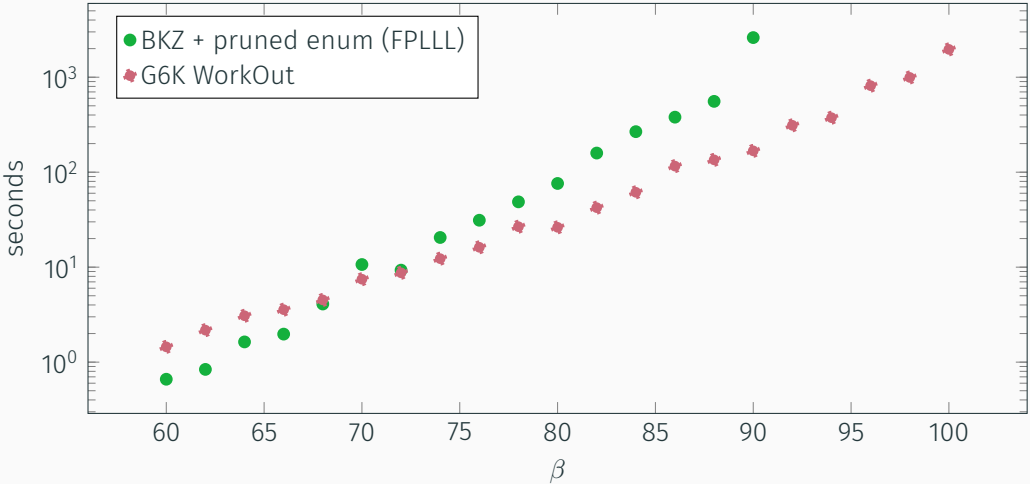

Average time in seconds for solving exact SVP

Darmstadt $\mathrm{HSVP}_{1.05}$ Challenges

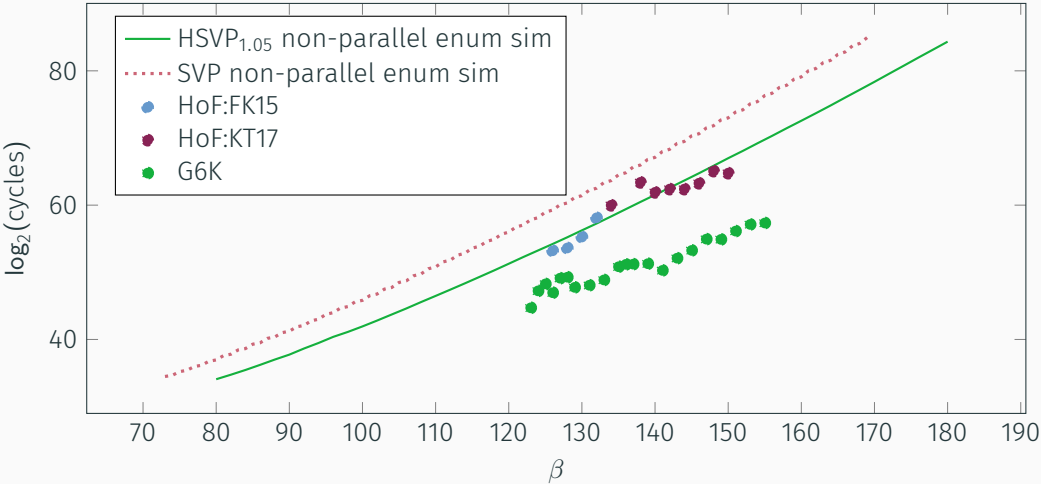

Estimated and reported costs for solving Darmstadt SVP Challenges.

Sieving: Open Questions

- G6K does not support coarse grained parallelism across different machines yet: not clear how exponential memory requirement scales in this regime
- Practical performance of asymptotically faster sieves still unclear
- Dedicated hardware ...

Quantum Estimates

Type	Cost Model \ Scheme	Kyber	NewHope	NTRU HRSS	SNTRU'
classical	$0.292\,\beta$	180	259	136	155
quantum	$0.265\,\beta$	163	235	123	140
classical	$1/(2e)\,\beta\log(\beta) - \beta + 16.1$	456	738	313	370
quantum	$1/(4e)\,\beta\log(\beta) - 1/2\beta + 8$	228	369	157	187

Sieving Given some vector $\mathbf{w}$ and a list of vectors L, apply Grover's algorithm to find $\{\mathbf{v} \in L \text{ s.t. } \|\mathbf{v} \pm \mathbf{w}\| \leq \|\mathbf{w}\|\}$.[13]

Enumeration Apply Montanaro's quantum backtracking algorithm for quadratic speed-up.[14]

[13] Thijs Laarhoven. Search problems in cryptography: From fingerprinting to lattice sieving. PhD thesis. Eindhoven University of Technology, 2015.

[14] Yoshinori Aono, Phong Q. Nguyen, and Yixin Shen. Quantum Lattice Enumeration and Tweaking Discrete Pruning. Cryptology ePrint Archive, Report 2018/546. `https://eprint.iacr.org/2018/546`. 2018.

Quantum Sieving

- A quantum sieve needs list of $2^{0.2075\beta}$ vectors before pairwise search with Grover
- Newer sieves use that the search is structured, Grover does unstructured search
  - Quantum Gauss Sieve

$$2^{(0.2075+\frac{1}{2}0.2075)\,\beta+o(\beta)} = 2^{0.311\,\beta+o(\beta)} \text{ time}, \qquad 2^{0.2075\,\beta+o(\beta)} \text{ memory}$$

  - Classical BGJ Sieve[15]

$$2^{0.311\,\beta+o(\beta)} \text{ time}, \qquad 2^{0.2075\,\beta+o(\beta)} \text{ memory}$$

- Asymptotically fastest sieves have small lists and thus less Grover speed-up potential

[15] Anja Becker, Nicolas Gama, and Antoine Joux. Speeding-up lattice sieving without increasing the memory, using sub-quadratic nearest neighbor search. Cryptology ePrint Archive, Report 2015/522. `http://eprint.iacr.org/2015/522`. 2015.

Implementing Quantum Algorithms for SVP

- The major operation in a sieve is to check whether two vectors reduce to some smaller vector.
- This can be implemented using the XOR and popcount trick $\Rightarrow$ the quantum circuit is relatively small.
- Sieving requires exponentially large quantum accessible RAM (qRAM). It is not clear that this can be build efficiently (due to error correction being required).
- Enumeration requires relatively high precision floating point arithmetic.
- Thus, quantum circuit for enumeration is likely to be larger than for sieving.
- But no exponential qRAM.

A Word on Lower Bounds

Cost Model \ Scheme	Kyber	NewHope	NTRU HRSS	SNTRU'
$0.292\,\beta$	180	259	136	155
$1/(2e)\,\beta\log(\beta) - \beta + 16.1$	456	738	313	370
$1/8\,\beta\log(\beta) - 0.75\beta + 2.3$	248	416	165	200
$0.265\,\beta$	163	235	123	140
$1/(4e)\,\beta\log(\beta) - 1/2\beta + 8$	228	369	157	187

These estimates ignore:

- (large) polynomial factors hidden in $o(\beta)$
- MAXDEPTH of quantum computers
- cost of a Grover iteration

Thus:

- cannot claim parameters need to be adjusted when these estimates are lowered
- must be careful about conclusions drawn in these models: some attacks don't work here but work in reality

Alternative Approaches

Dual Find short $\mathbf{v}$ s.t. $\mathbf{v} \cdot \mathbf{A}$ is short.

BKW combinatorial technique, relatively efficient for small secrets

Arora-Ge use Gröbner bases, asymptotically efficient , but large constants in the exponent

Hybrid Attack combine combinatorial techniques with lattice reduction

Rule of Thumb

Don't need to worry about last three unless secret is unusually small (e.g. ternary) and/or sparse.

More Open Questions

- Many submissions use small and sparse secrets where combinatorial techniques apply. Cost of these not fully understood.
- (Structured) Ideal-SVP is easier than General SVP on a quantum computer.[16] Ring-LWE (but for a choice of parameters typically not used in practice) is at least as hard as Ideal-SVP, but it is not clear if it is harder, e.g. if those attacks extend.
- The effect of decryption failures in probabilistic encryption based on LWE not fully understood. Some submissions completely eliminate these.

[16] Ronald Cramer, Léo Ducas, and Benjamin Wesolowski. Short Stickelberger Class Relations and Application to Ideal-SVP. In: *EUROCRYPT 2017, Part I*. ed. by Jean-Sébastien Coron and Jesper Buus Nielsen. Vol. 10210. LNCS. Springer, Heidelberg, 2017, pp. 324–348. DOI: 10.1007/978-3-319-56620-7_12.

Thank You

References I

[Alb+17] Martin R. Albrecht, Florian Göpfert, Fernando Virdia, and Thomas Wunderer. Revisiting the Expected Cost of Solving uSVP and Applications to LWE. In: *ASIACRYPT 2017, Part I*. Ed. by Tsuyoshi Takagi and Thomas Peyrin. Vol. 10624. LNCS. Springer, Heidelberg, Dec. 2017, pp. 297–322. DOI: 10.1007/978-3-319-70694-8_11.

[Alb+18] Martin R. Albrecht, Benjamin R. Curtis, Amit Deo, Alex Davidson, Rachel Player, Eamonn W. Postlethwaite, Fernando Virdia, and Thomas Wunderer. Estimate All the LWE, NTRU Schemes! In: *SCN 18*. Ed. by Dario Catalano and Roberto De Prisco. Vol. 11035. LNCS. Springer, Heidelberg, Sept. 2018, pp. 351–367. DOI: 10.1007/978-3-319-98113-0_19.

[Alb+19] Martin R. Albrecht, Léo Ducas, Gottfried Herold, Elena Kirshanova, Eamonn W. Postlethwaite, and Marc Stevens. The General Sieve Kernel and New Records in Lattice Reduction. In: *EUROCRYPT 2019, Part II*. Ed. by Yuval Ishai and Vincent Rijmen. Vol. 11477. LNCS. Springer, Heidelberg, May 2019, pp. 717–746. DOI: 10.1007/978-3-030-17656-3_25.

[Alk+16] Erdem Alkim, Léo Ducas, Thomas Pöppelmann, and Peter Schwabe. Post-quantum Key Exchange - A New Hope. In: *USENIX Security 2016*. Ed. by Thorsten Holz and Stefan Savage. USENIX Association, Aug. 2016, pp. 327–343.

[ANS18] Yoshinori Aono, Phong Q. Nguyen, and Yixin Shen. Quantum Lattice Enumeration and Tweaking Discrete Pruning. Cryptology ePrint Archive, Report 2018/546. https://eprint.iacr.org/2018/546. 2018.

[APS15] Martin R. Albrecht, Rachel Player, and Sam Scott. On the concrete hardness of Learning with Errors. In: *Journal of Mathematical Cryptology* 9.3 (2015), pp. 169–203.

References II

[Bec+16] Anja Becker, Léo Ducas, Nicolas Gama, and Thijs Laarhoven. New directions in nearest neighbor searching with applications to lattice sieving. In: *27th SODA*. Ed. by Robert Krauthgamer. ACM-SIAM, Jan. 2016, pp. 10–24. DOI: 10.1137/1.9781611974331.ch2.

[BGJ15] Anja Becker, Nicolas Gama, and Antoine Joux. Speeding-up lattice sieving without increasing the memory, using sub-quadratic nearest neighbor search. Cryptology ePrint Archive, Report 2015/522. http://eprint.iacr.org/2015/522. 2015.

[CDW17] Ronald Cramer, Léo Ducas, and Benjamin Wesolowski. Short Stickelberger Class Relations and Application to Ideal-SVP. In: *EUROCRYPT 2017, Part I*. Ed. by Jean-Sébastien Coron and Jesper Buus Nielsen. Vol. 10210. LNCS. Springer, Heidelberg, 2017, pp. 324–348. DOI: 10.1007/978-3-319-56620-7_12.

[Che13] Yuanmi Chen. Réduction de réseau et sécurité concrète du chiffrement complètement homomorphe. PhD thesis. Paris 7, 2013.

[HPS11] Guillaume Hanrot, Xavier Pujol, and Damien Stehlé. Analyzing Blockwise Lattice Algorithms Using Dynamical Systems. In: *CRYPTO 2011*. Ed. by Phillip Rogaway. Vol. 6841. LNCS. Springer, Heidelberg, Aug. 2011, pp. 447–464. DOI: 10.1007/978-3-642-22792-9_25.

[HS07] Guillaume Hanrot and Damien Stehlé. Improved Analysis of Kannan's Shortest Lattice Vector Algorithm. In: *CRYPTO 2007*. Ed. by Alfred Menezes. Vol. 4622. LNCS. Springer, Heidelberg, Aug. 2007, pp. 170–186. DOI: 10.1007/978-3-540-74143-5_10.

References III

[Laa15] Thijs Laarhoven. Search problems in cryptography: From fingerprinting to lattice sieving. PhD thesis. Eindhoven University of Technology, 2015.

[MV10] Daniele Micciancio and Panagiotis Voulgaris. Faster Exponential Time Algorithms for the Shortest Vector Problem. In: *21st SODA*. Ed. by Moses Charika. ACM-SIAM, Jan. 2010, pp. 1468–1480. DOI: 10.1137/1.9781611973075.119.

[Ngu10] Phong Q. Nguyen. Hermite's Constant and Lattice Algorithms. In: ed. by Phong Q. Nguyen and Brigitte Vallée. ISC. Springer, Heidelberg, 2010, pp. 19–69. ISBN: 978-3-642-02294-4. DOI: 10.1007/978-3-642-02295-1.

[Pho+17] Le Trieu Phong, Takuya Hayashi, Yoshinori Aono, and Shiho Moriai. LOTUS. Tech. rep. available at https://csrc.nist.gov/projects/post-quantum-cryptography/round-1-submissions. National Institute of Standards and Technology, 2017.